

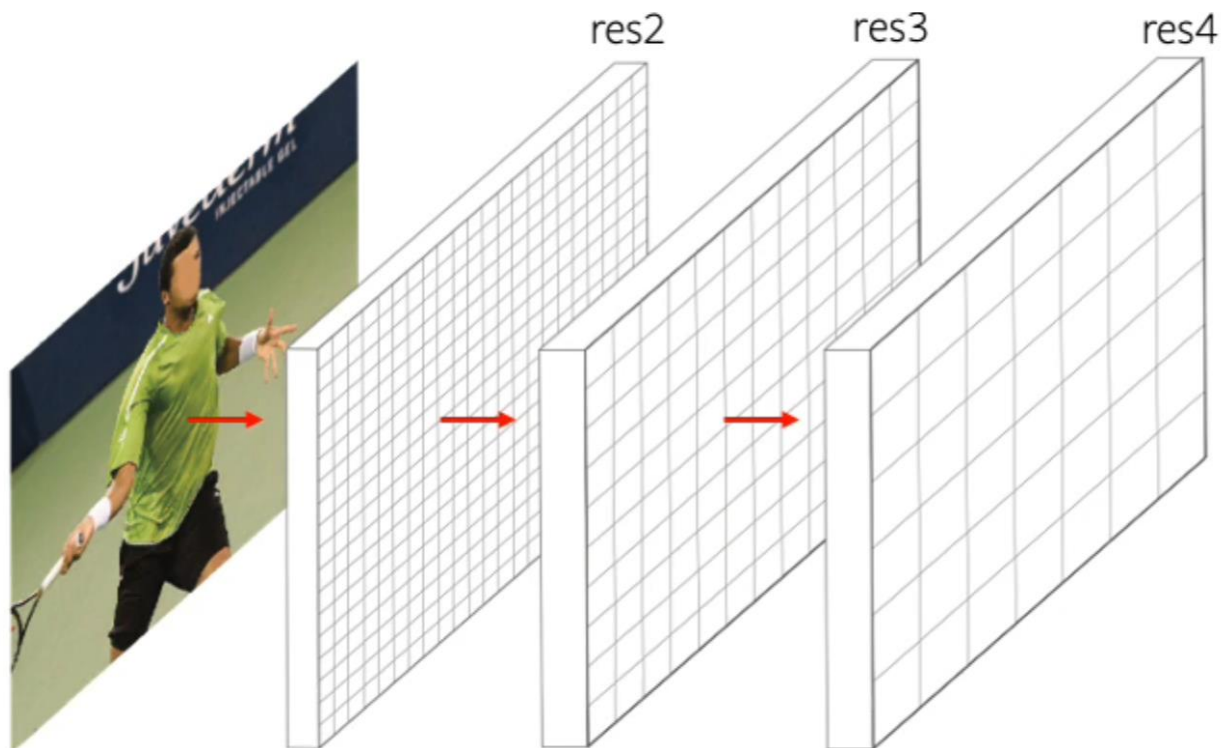
# PointRend

Image Segmentation as  
Rendering

**Su Hyung Choi**

Korea University

Research Intern @ Computer Vision Lab



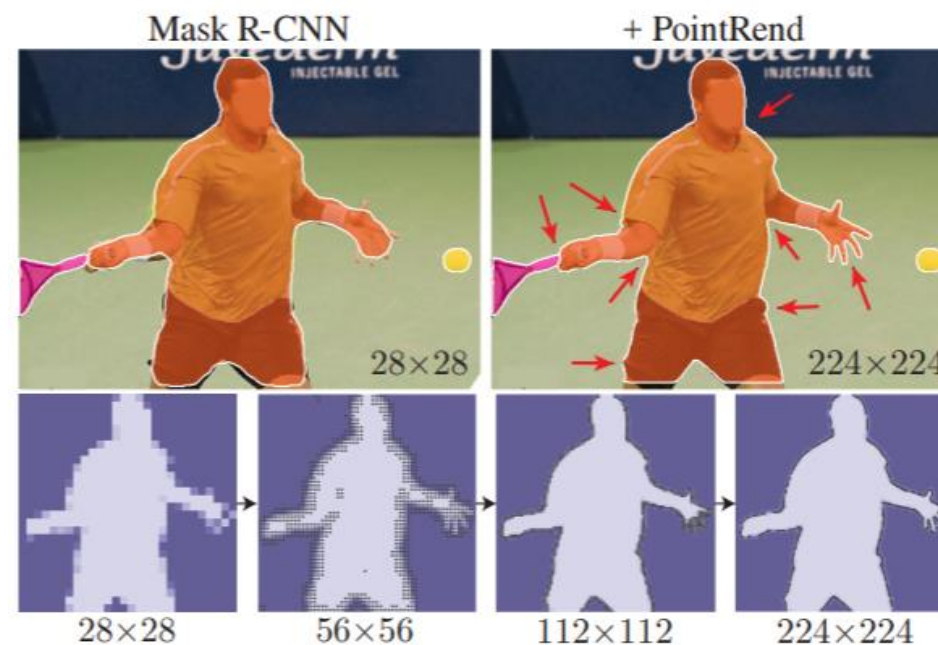
**Submitted on Jan 11, 2022**

Authors: Facebook AI Research (FAIR)

Alexander Kirillov, Yuxin Wu, Kaiming He Ross Girshick

# Typical Problems in Image Segmentation

- CNNs for image segmentation typically operate on regular 3 channel grids.
  - Not computationally ideal for image segmentation.
  - Label maps predicted by these networks are mostly smooth.
  - Unnecessarily oversample the smooth areas while simultaneously under-sampling object boundaries.
- The result is excess computation in smooth regions and blurry contours.
- Often predict labels on a low-resolution regular grid.



# PointRend

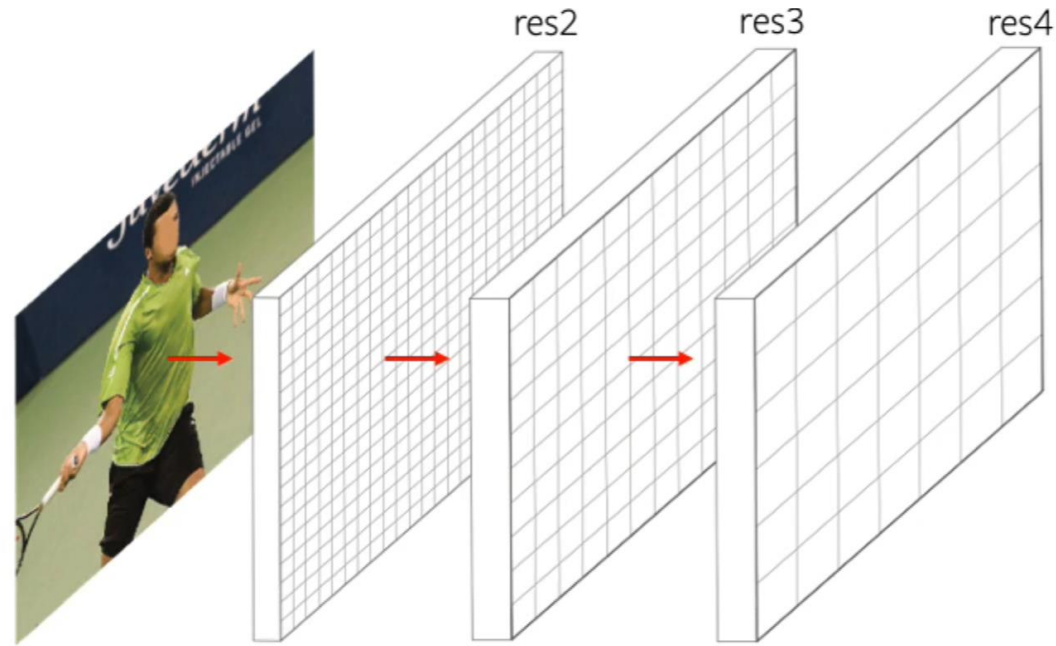
Propose “Point based Rendering” as a methodology for image segmentation using point representations.

To Efficiently “render” high quality label map

Benefits over other architectures:

1. Training and Inference is less computationally intensive
2. Generate higher quality, higher resolution label maps (masks)


# CNN/FCN is the main tool for 2D image recognition



- Operates with **feature map on regular grids**



Output resolution is tradeoff between computational cost and level of detail



7x7

14x14


28x28




56x56

112x112

224x224



Mask R-CNN efficiently predicts  
low-resolution masks



7x7

14x14

**28x28**




56x56

112x112

224x224



Note that different areas require different levels of detail



7x7

14x14

28x28



56x56

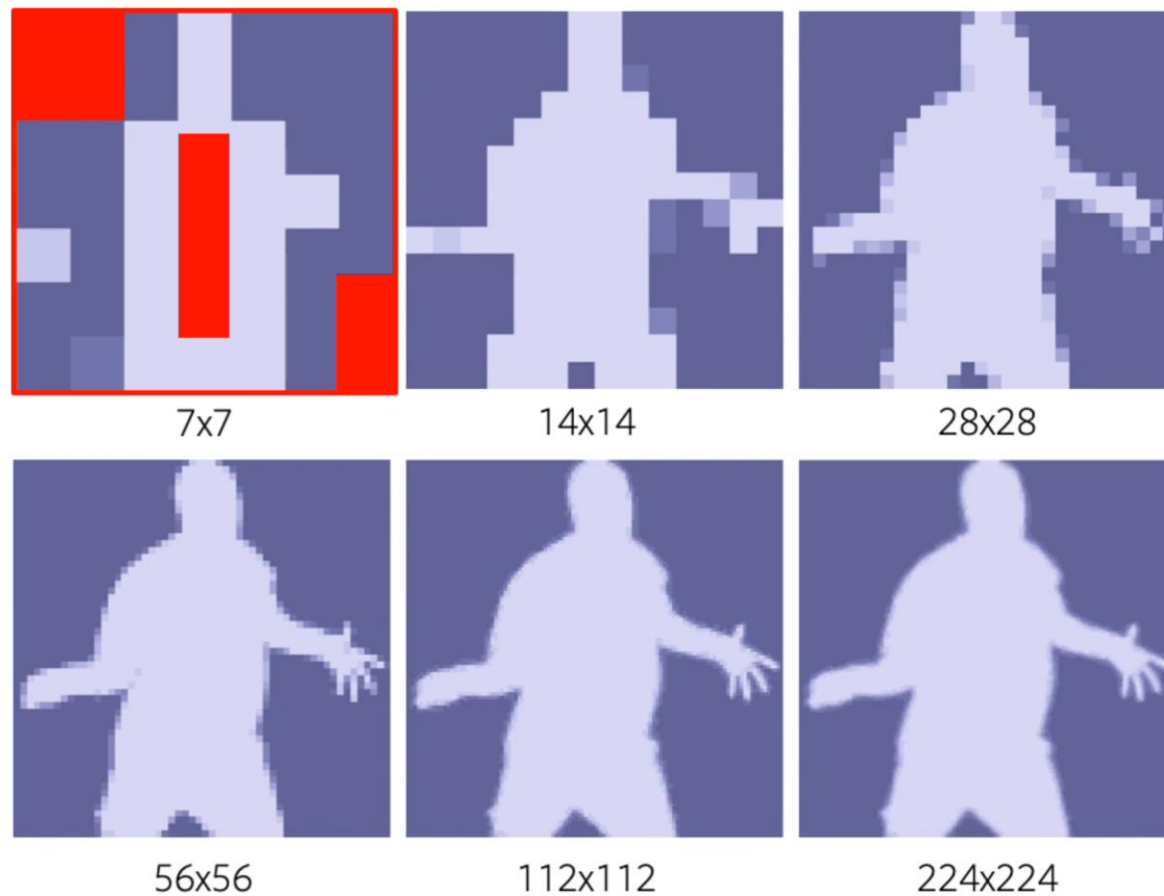
112x112

224x224






Some are perfectly segmented with  
low resolution prediction





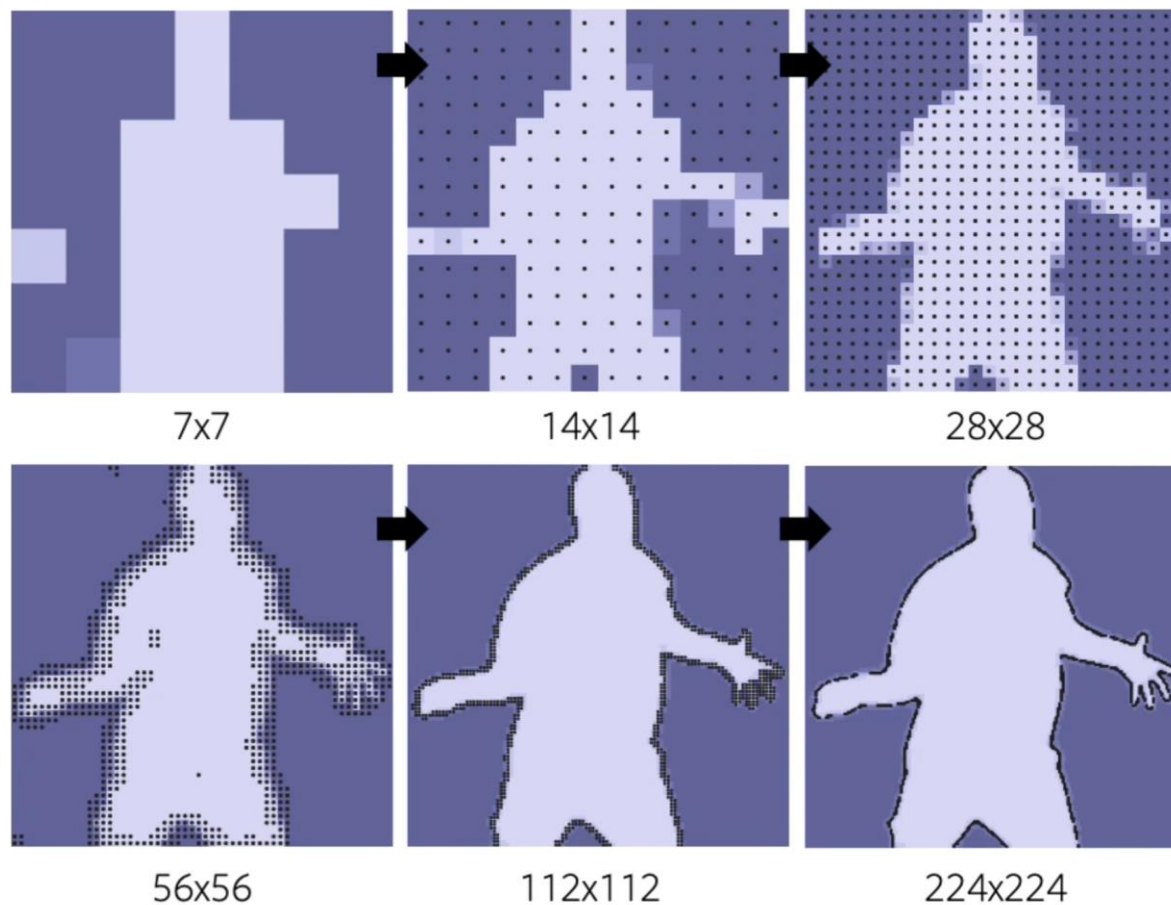


Whereas high-frequency regions require prediction in a very high resolution

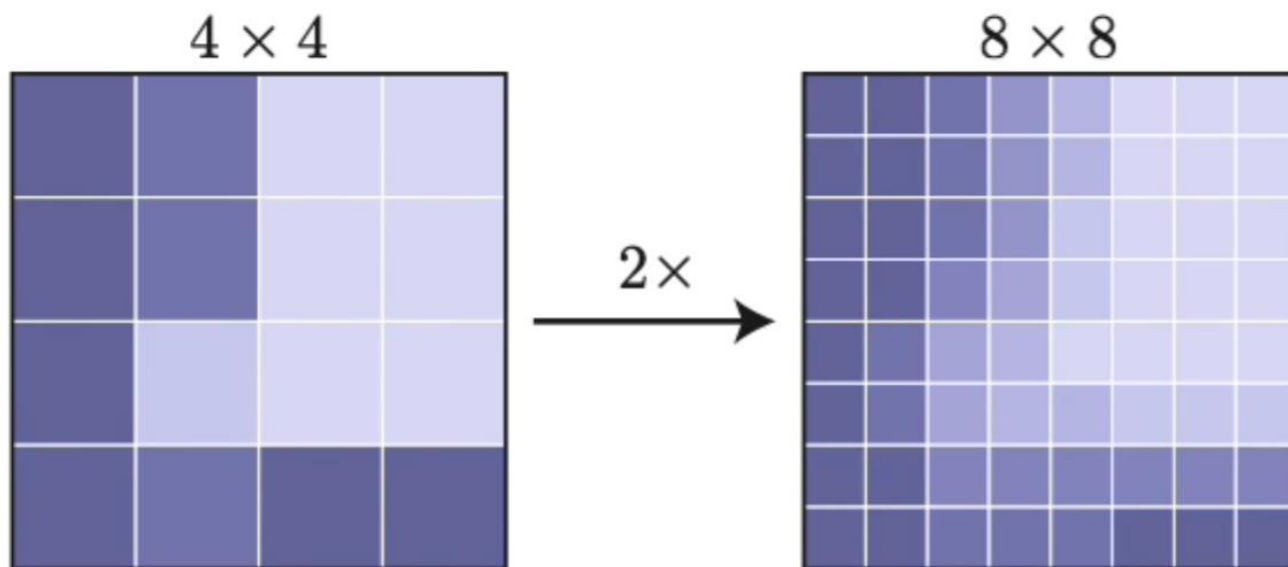




PointRend gradually increases resolution by making predictions for the most uncertain points

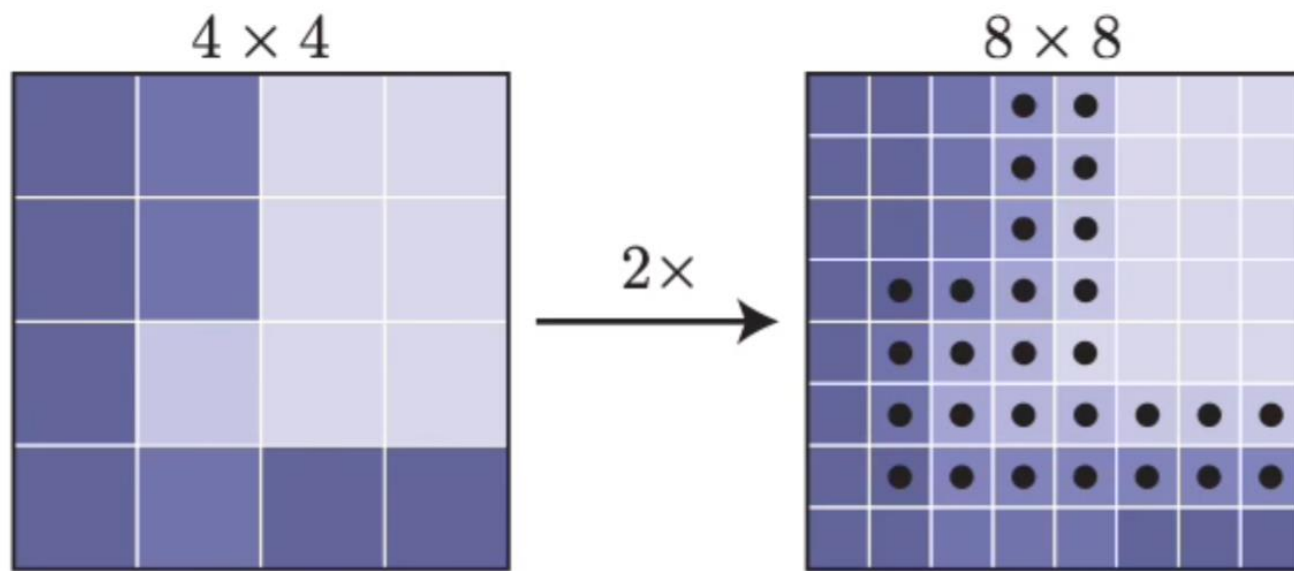


# Point-based inference via subdivision



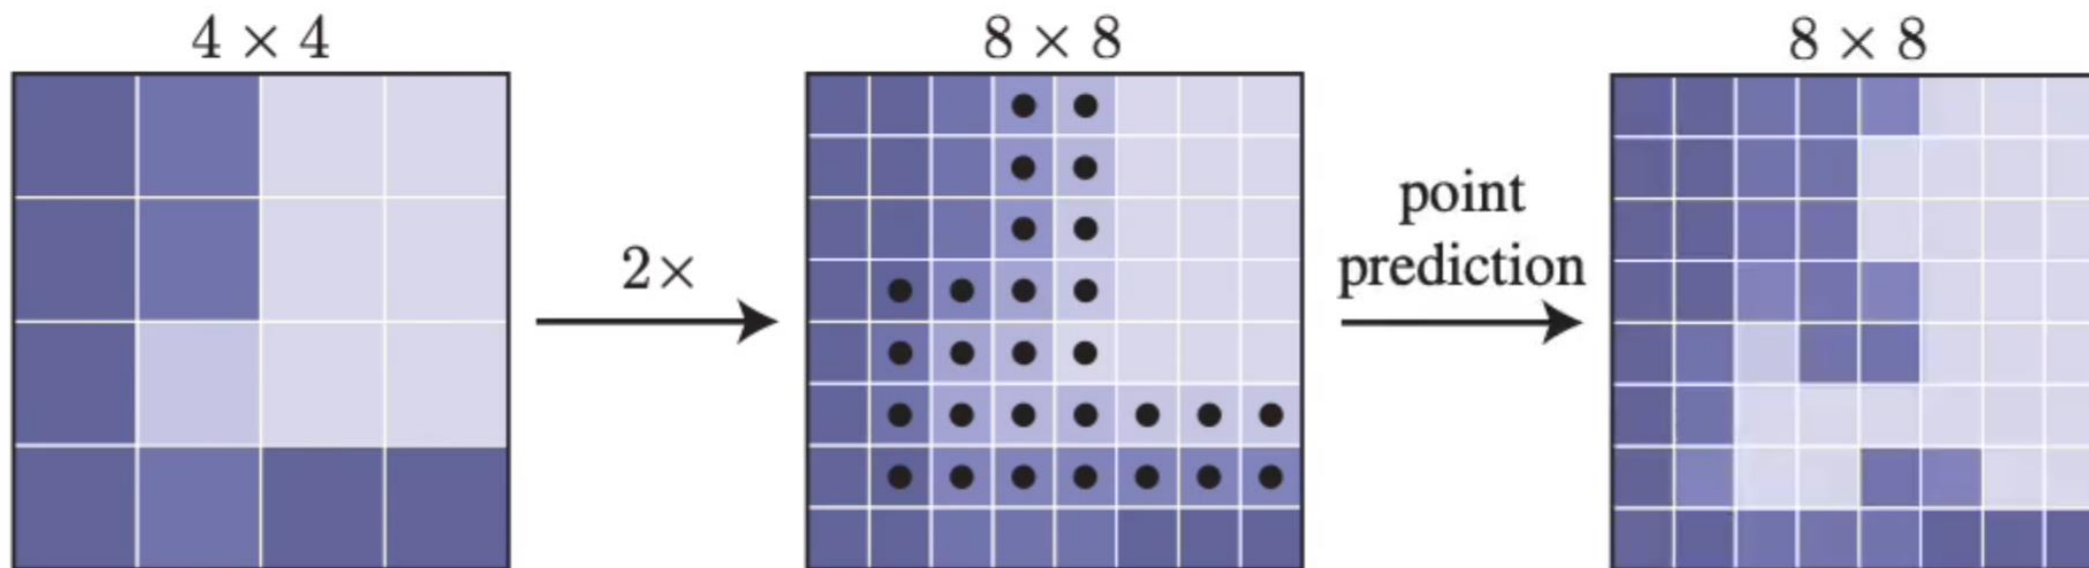
Lower resolutions prediction is  
unsampled with bilinear interpolation

# Point-based inference via subdivision



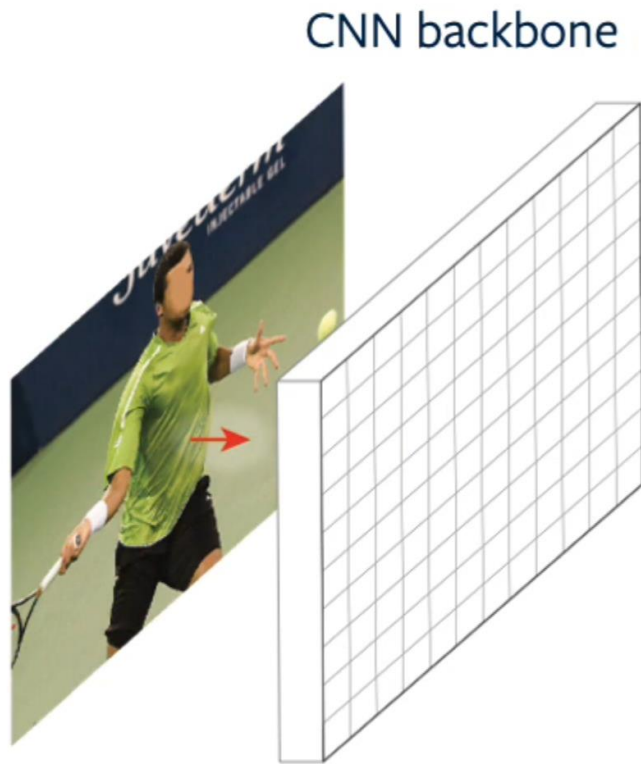
The subset of most uncertain points is selected

# Point-based inference via subdivision



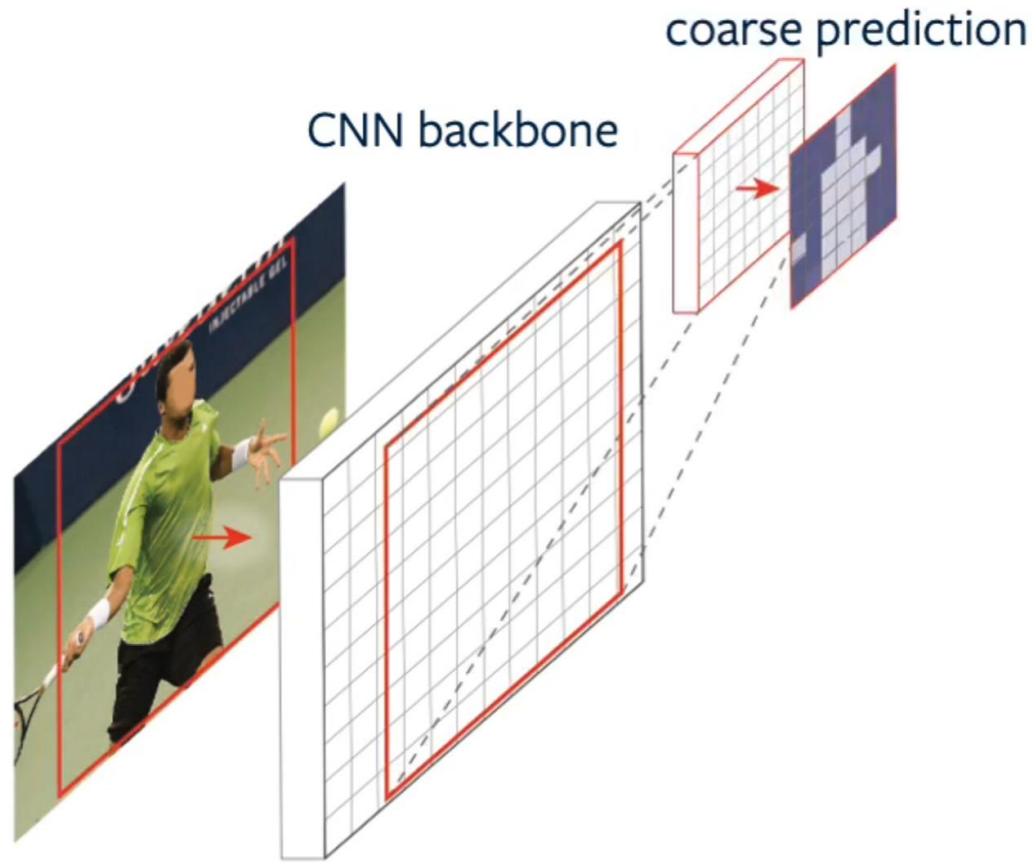
Prediction for each selected points is refined using a lightweight MLP

# PointRender architecture for instance segmentation



Backbone computes features  
that represent the whole image


# PointRender architecture for instance segmentation



For each bounding box  
a small head yields  
low-resolution mask prediction

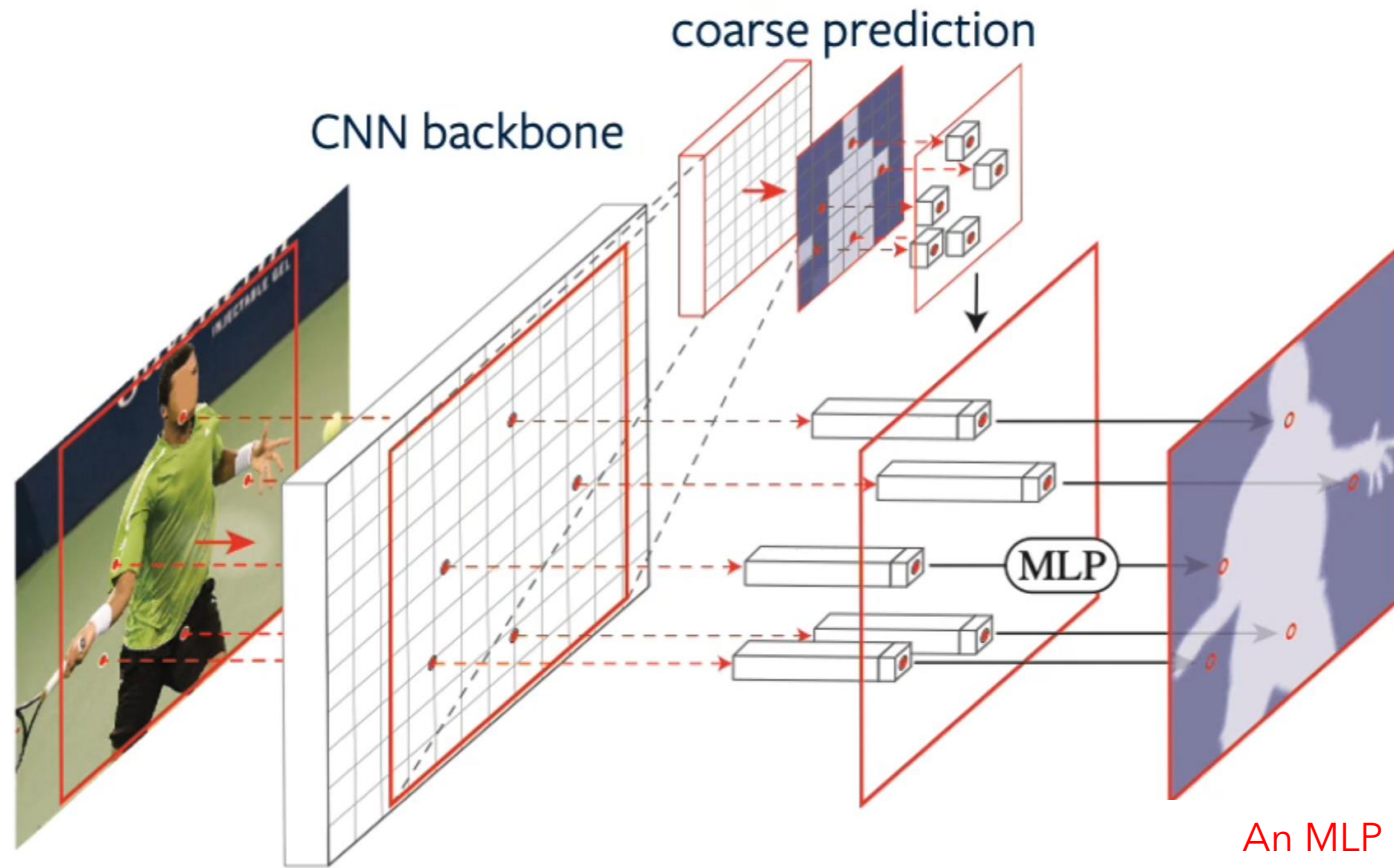


# PointRend architecture for instance segmentation



For a subset of points we extract features from the coarse prediction and the backbone features using bilinear interpolation

# PointRend architecture for instance segmentation

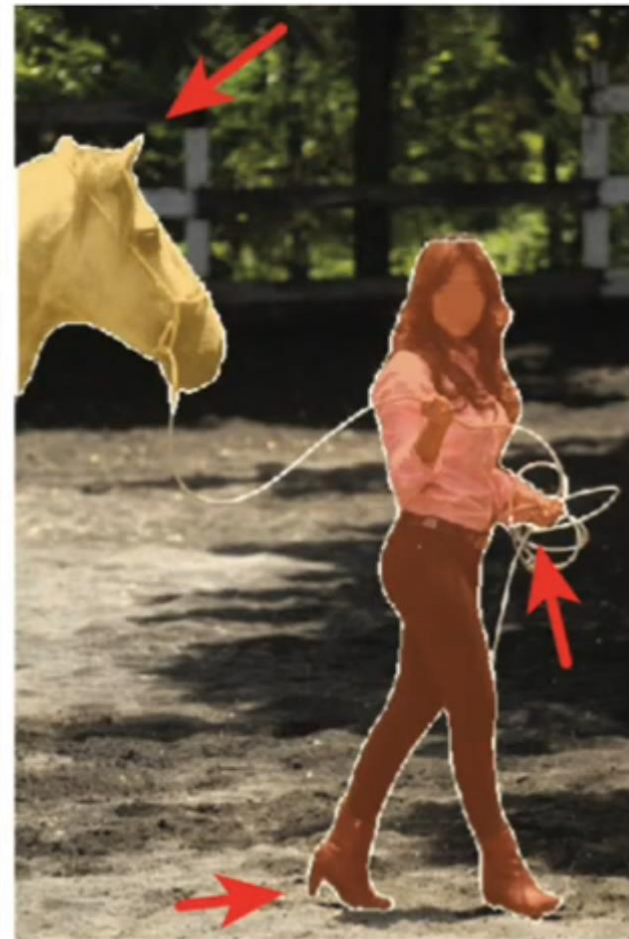
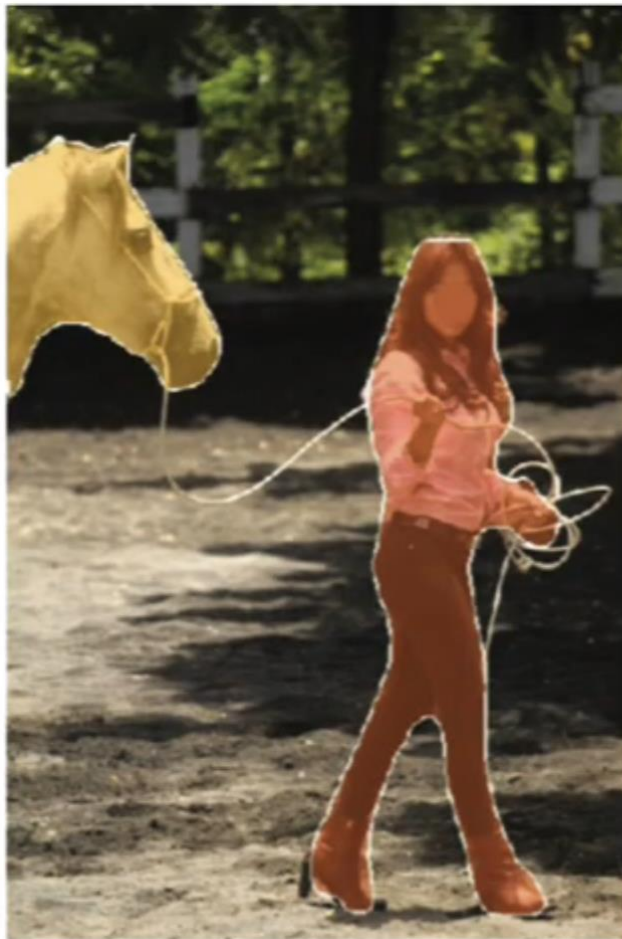


An MLP is used to make prediction for each point independently

# Mask R-CNN with standard head vs. Mask R-CNN with PointRend



**+ PointRend**



**+ PointRend**

# Quantitative comparison: instance segmentation

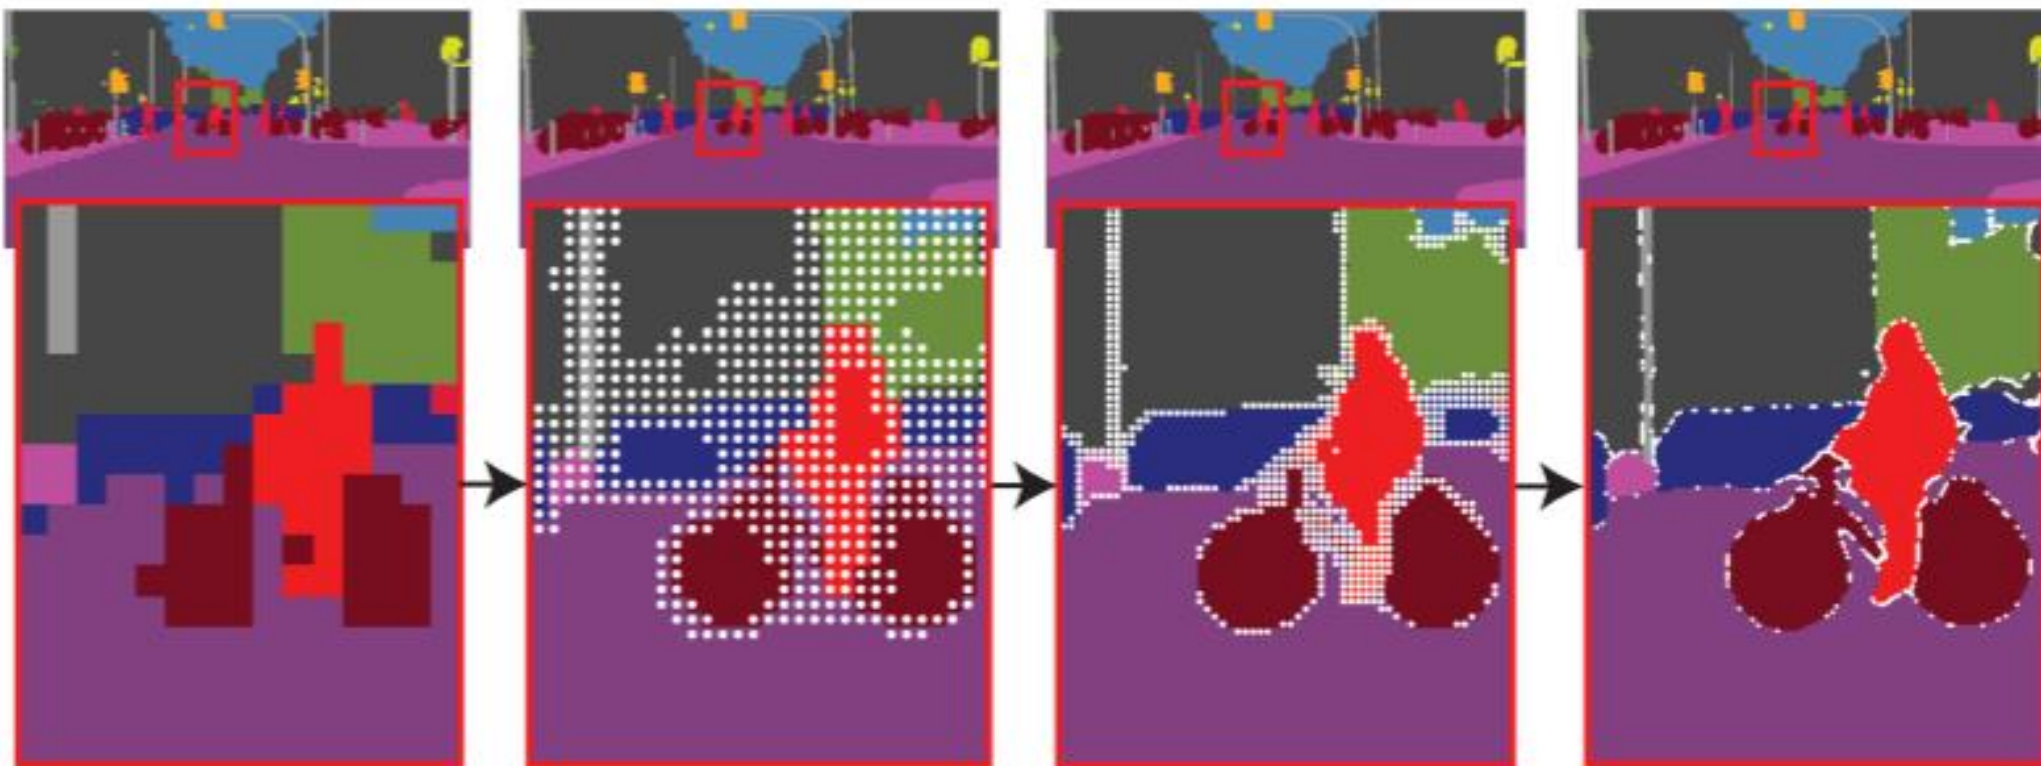
mask head	output resolution	COCO		Cityscapes AP
		AP	AP*	
4× conv	28×28	35.2	37.6	33.0
PointRend	224×224	<b>36.3</b> (+1.1)	<b>39.7</b> (+2.1)	<b>35.8</b> (+2.8)

Mask R-CNN with standard head (4x Conv) vs. Mask R-CNN with PointRend

AP\* is COCO mask AP for a COCO-trained model  
evaluated against the higher quality LVIS annotation

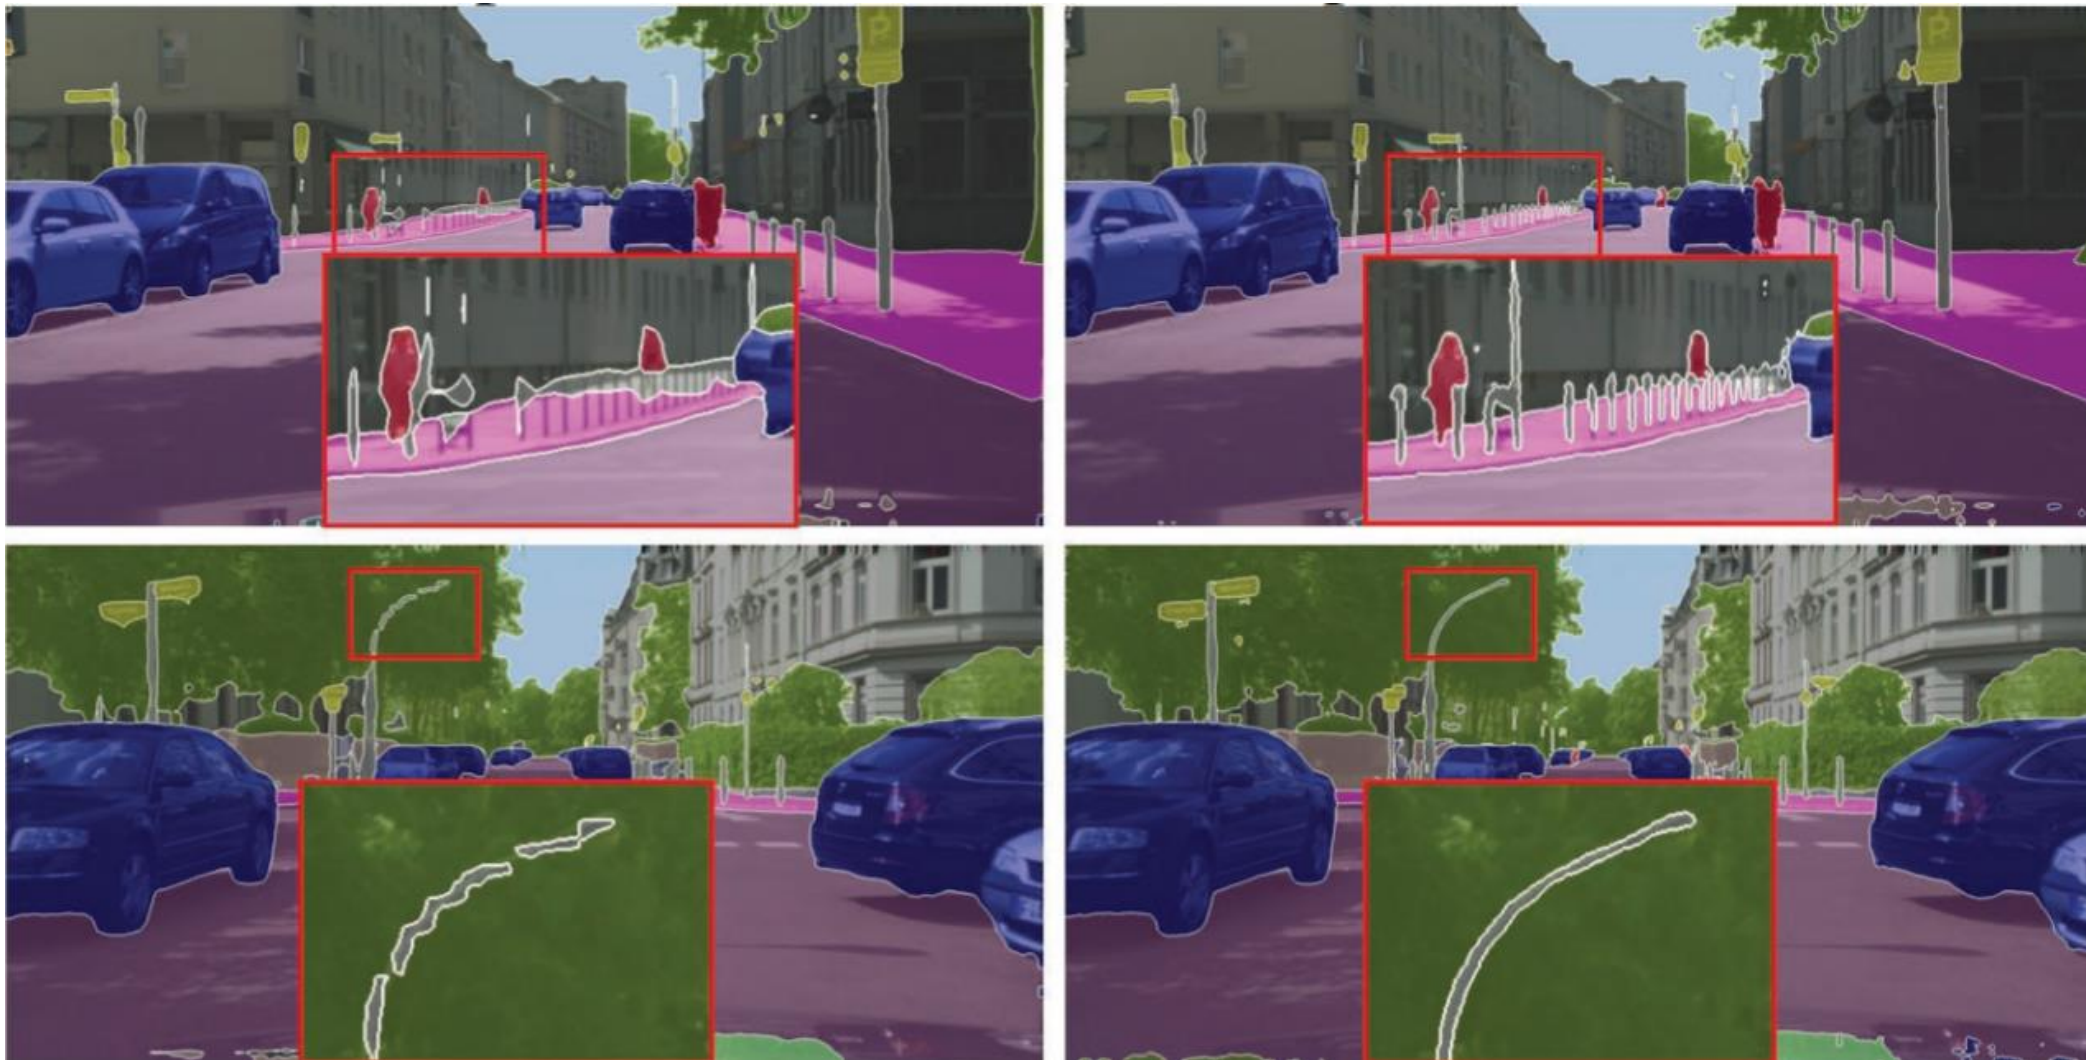


# PointRender for semantic segmentation



PointRender can be applied on top of  
any modern semantic segmentation model

# Deeplab V3 vs. Deeplab V3 + PointRend



**+ PointRend**

# Quantitative comparison: semantic segmentation

method	output resolution	mIoU
DeeplabV3-OS-16	$64 \times 128$	77.2
DeeplabV3-OS-8	$128 \times 256$	77.8 (+0.6)
DeeplabV3-OS-16 + PointRend	$1024 \times 2048$	<b>78.4</b> (+1.2)

## DeeplabV3 vs. DeeplabV3 with PointRend

method	output resolution	mIoU
SemanticFPN $P_2$ - $P_5$	$256 \times 512$	77.7
SemanticFPN $P_2$ - $P_5$ + PointRend	$1024 \times 2048$	<b>78.6</b> (+0.9)
SemanticFPN $P_3$ - $P_5$	$128 \times 256$	77.4
SemanticFPN $P_3$ - $P_5$ + PointRend	$1024 \times 2048$	<b>78.5</b> (+1.1)

## SemanticFPN vs SemanticFPN with PointRend



# Conclusion

- High resolution output with little to no computational overhead  
Higher resolution, more accurate masks  
with fewer model params, less compute time.
- “plug & play” on top of any FCN-based model for segmentation
- Significant quantitative and qualitative improvement

# Thank You.

[Paper Review] PointRend: Image Segmentation as Rendering

Su Hyung Choi