PointRend

Image Segmentation as Rendering

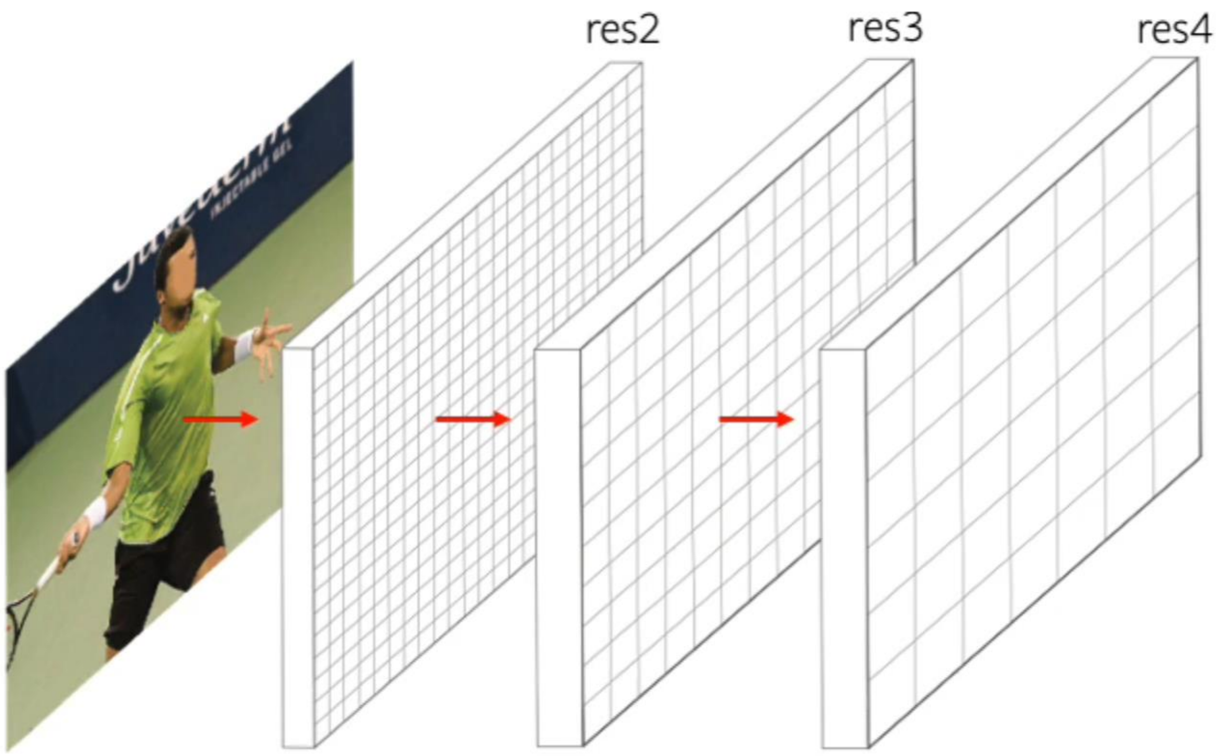

Su Hyung Choi

Korea University

Research Intern @ Computer Vision Lab

Submitted on Jan 6, 2022

Authors: Facebook AI Research (FAIR)

Alexander Kirillov, Yuxin Wu, Kaiming He Ross Girshick

Typical Problems in Image Segmentation

- CNNs for image segmentation typically operate on regular 3 channel grids.
- Not computationally ideal for image segmentation.
- Label maps predicted by these networks are mostly smooth.
- Unnecessarily oversample the smooth areas while simultaneously under-sampling object boundaries.
- The result is excess computation in smooth regions and blurry contours.
- Often predict labels on a low-resolution regular grid.

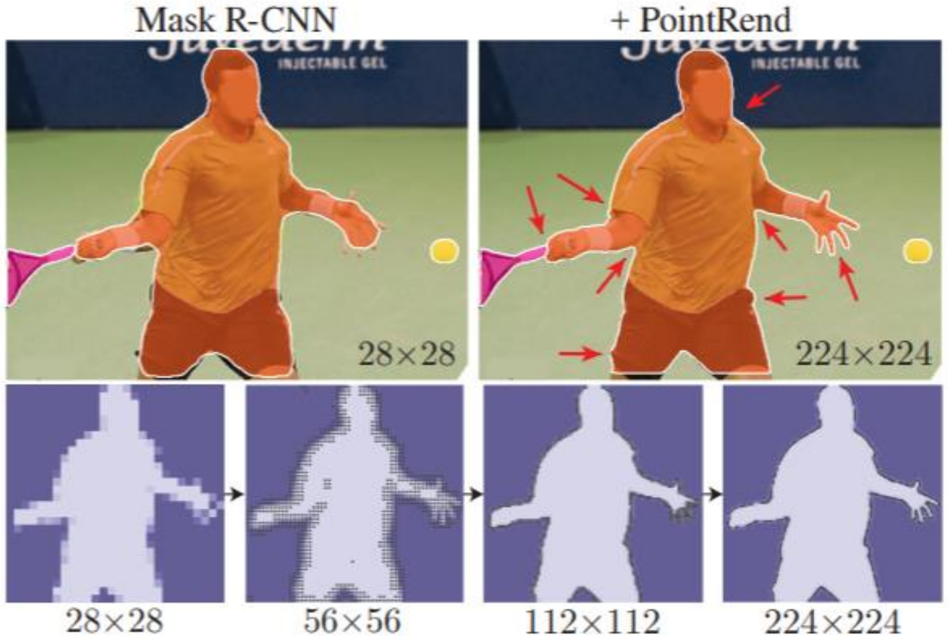

PointRend

Propose "Point based Rendering" as a methodology
for image segmentation using point representations.

To Efficiently "render" high quality label map

Benefits over other architectures:

1. Training and Inference is less computationally intensive
2. Generate higher quality, higher resolution label maps (masks)

CNN/FCN is the main tool for 2D image recognition

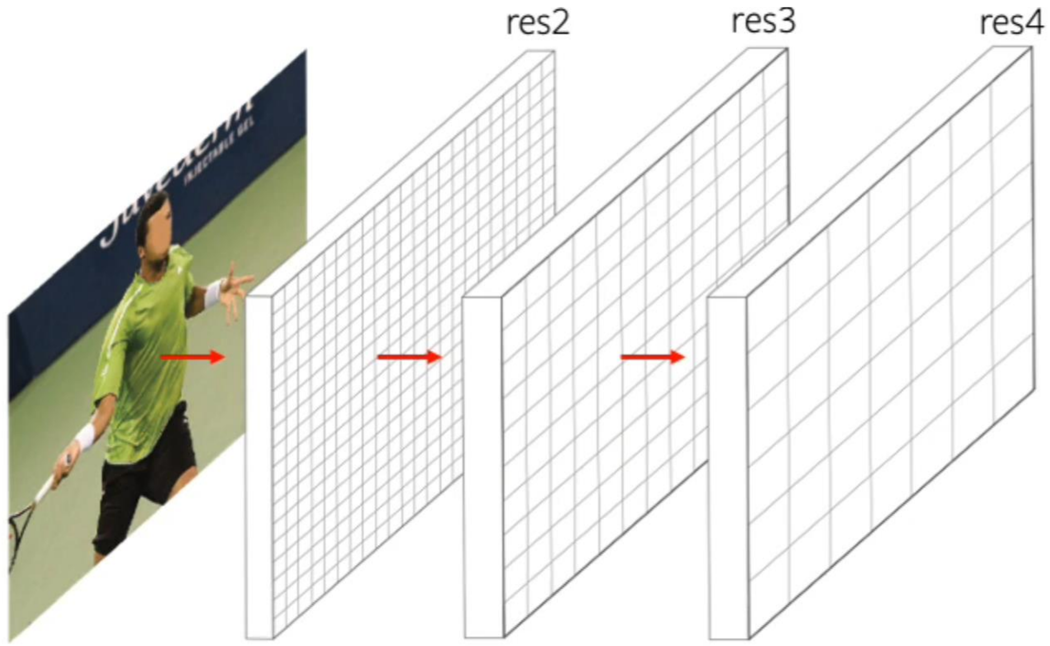

- Operates with **feature map on regular grids**

Output resolution is tradeoff between computational cost and level of detail

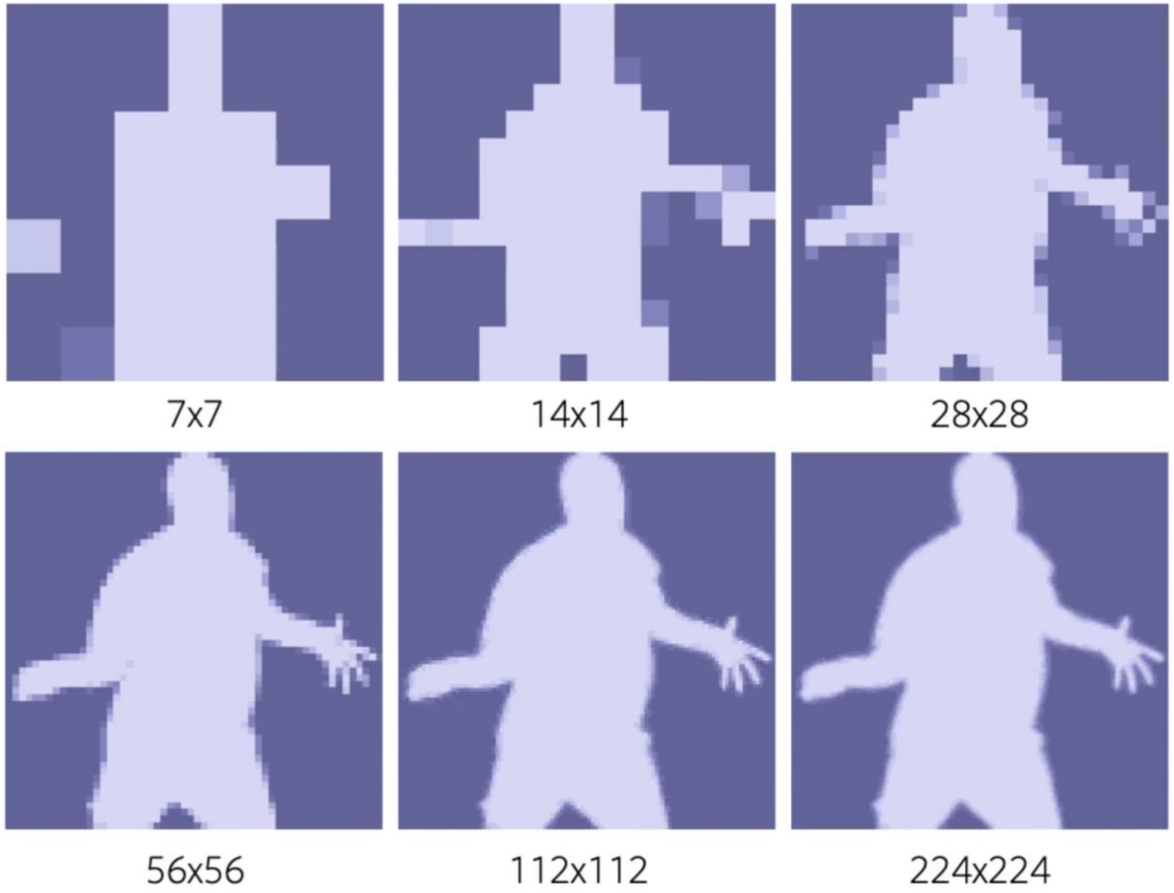

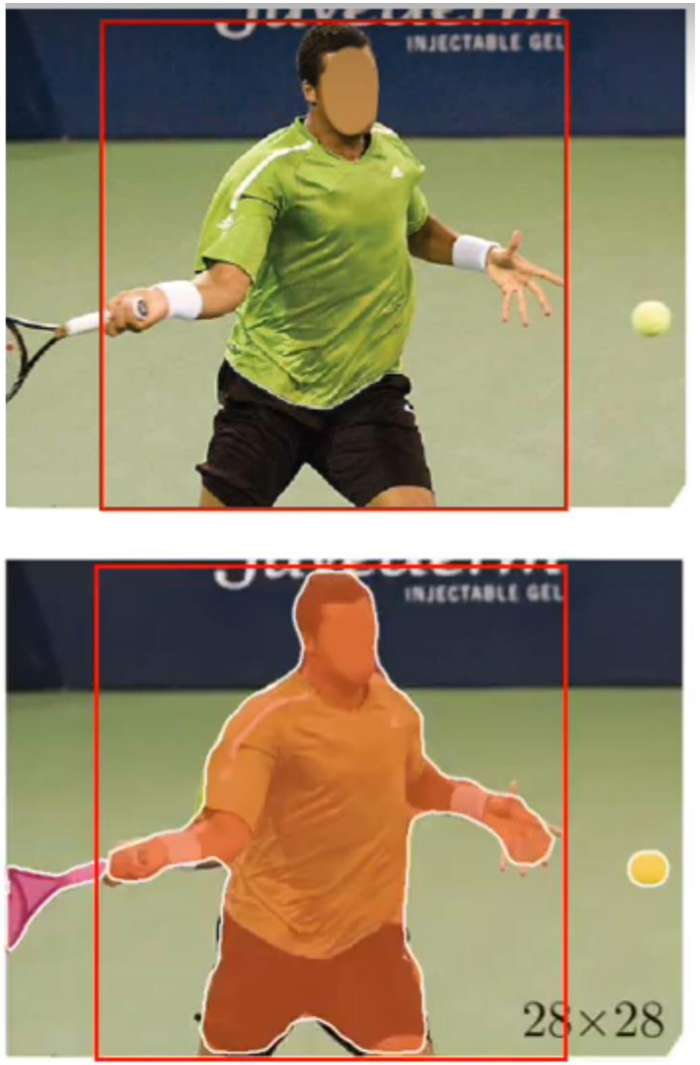

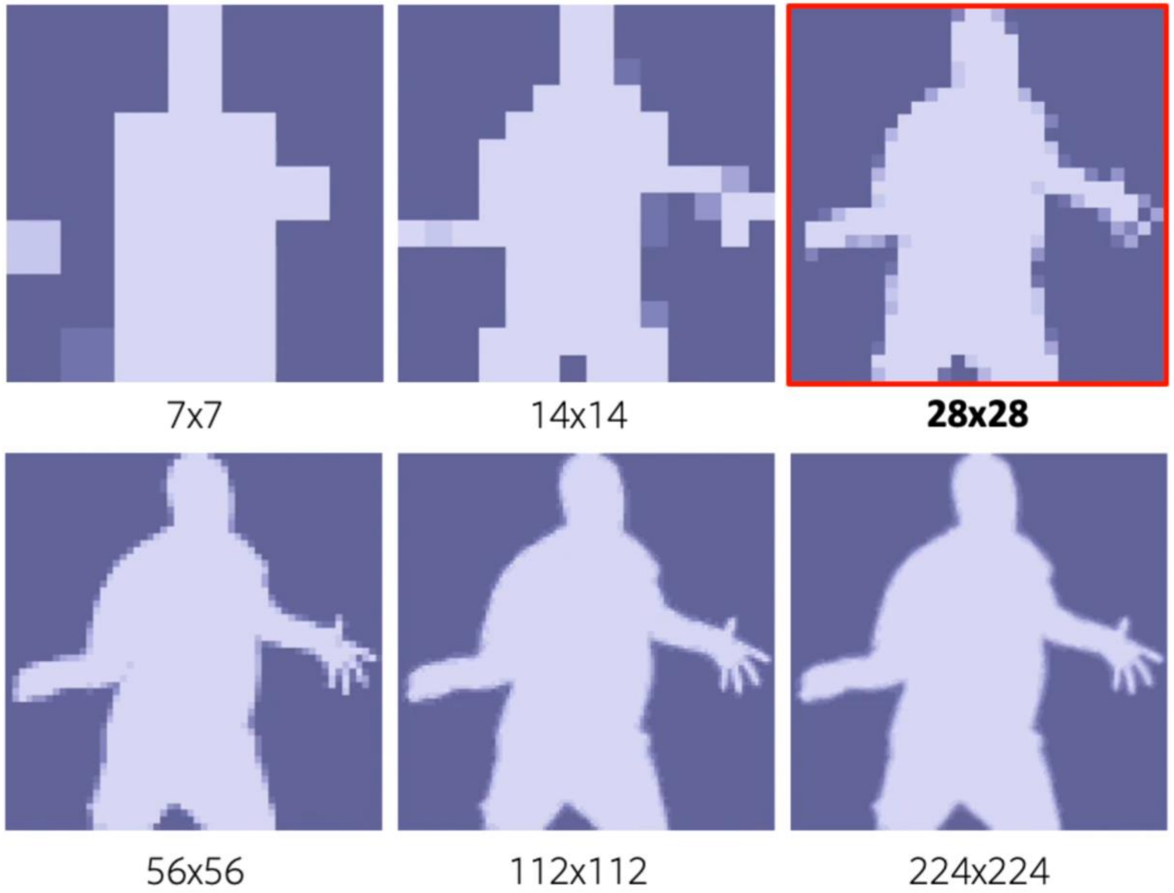

Mask R-CNN efficiently predicts
low-resolution masks

Note that different areas require different levels of detail

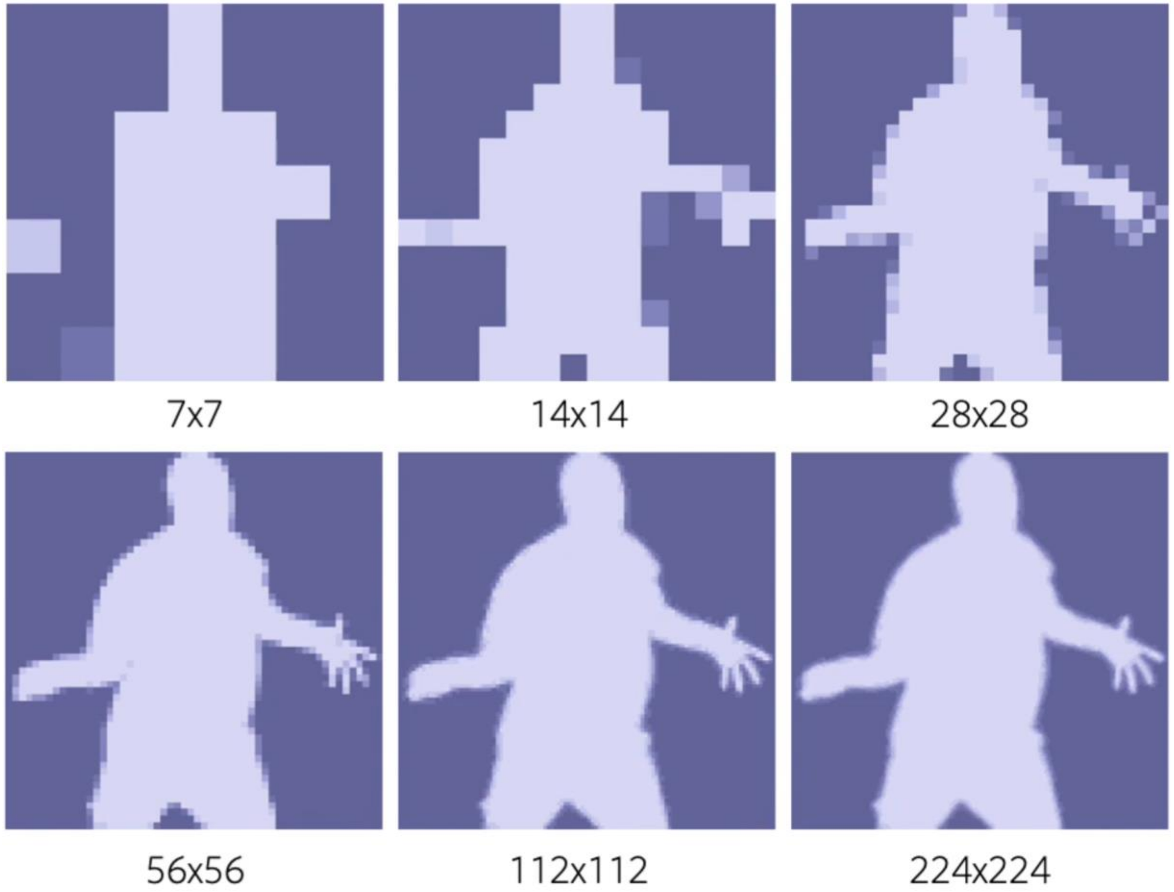

Some are perfectly segmented with low resolution prediction

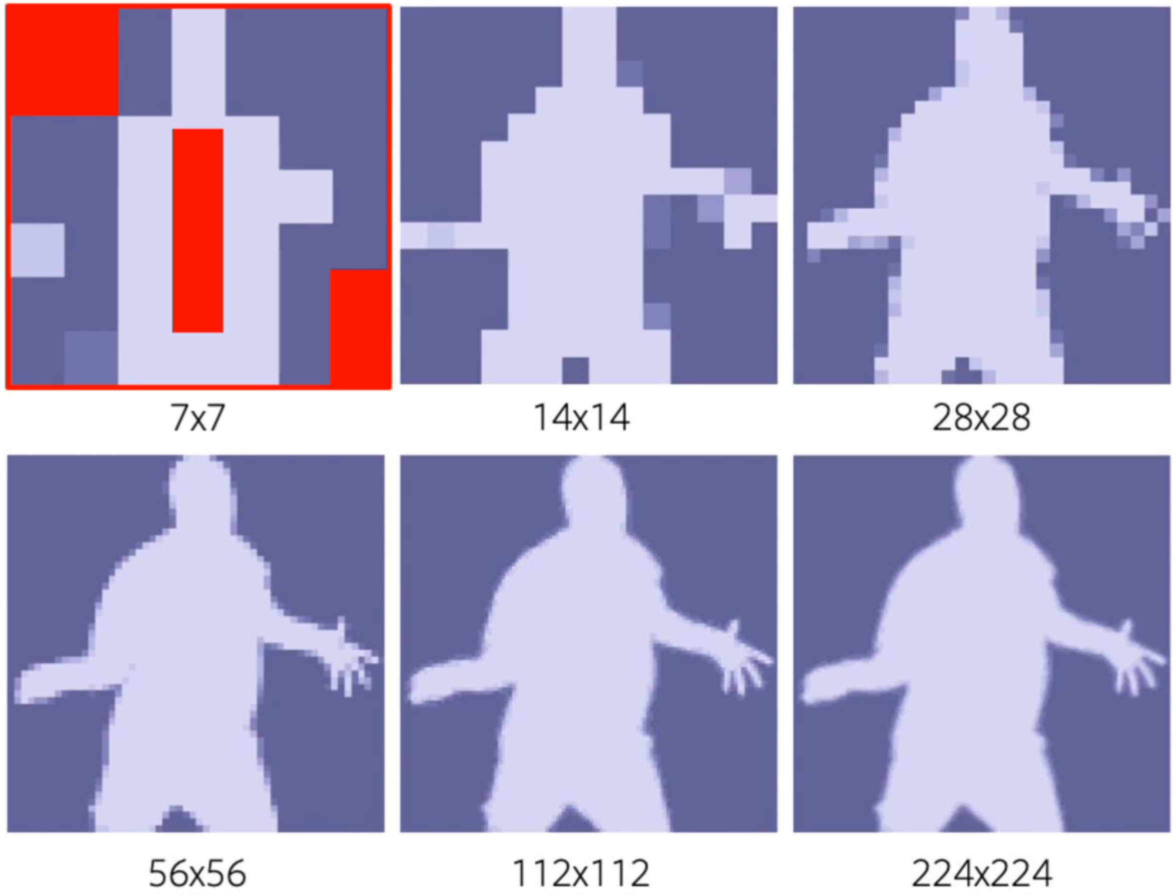

Whereas high-frequency regions require prediction in a very high resolution

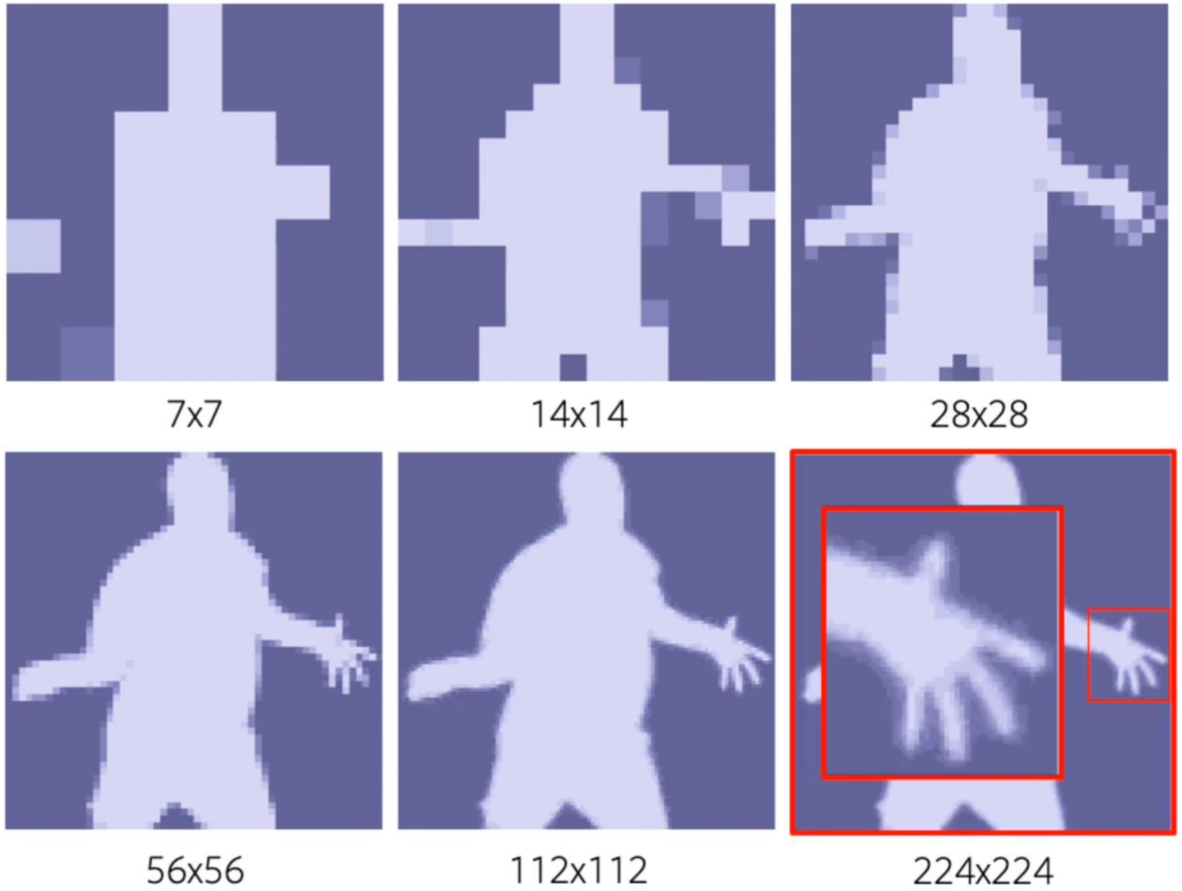

PointRend gradually increases
resolution by making predictions
for the most uncertain points

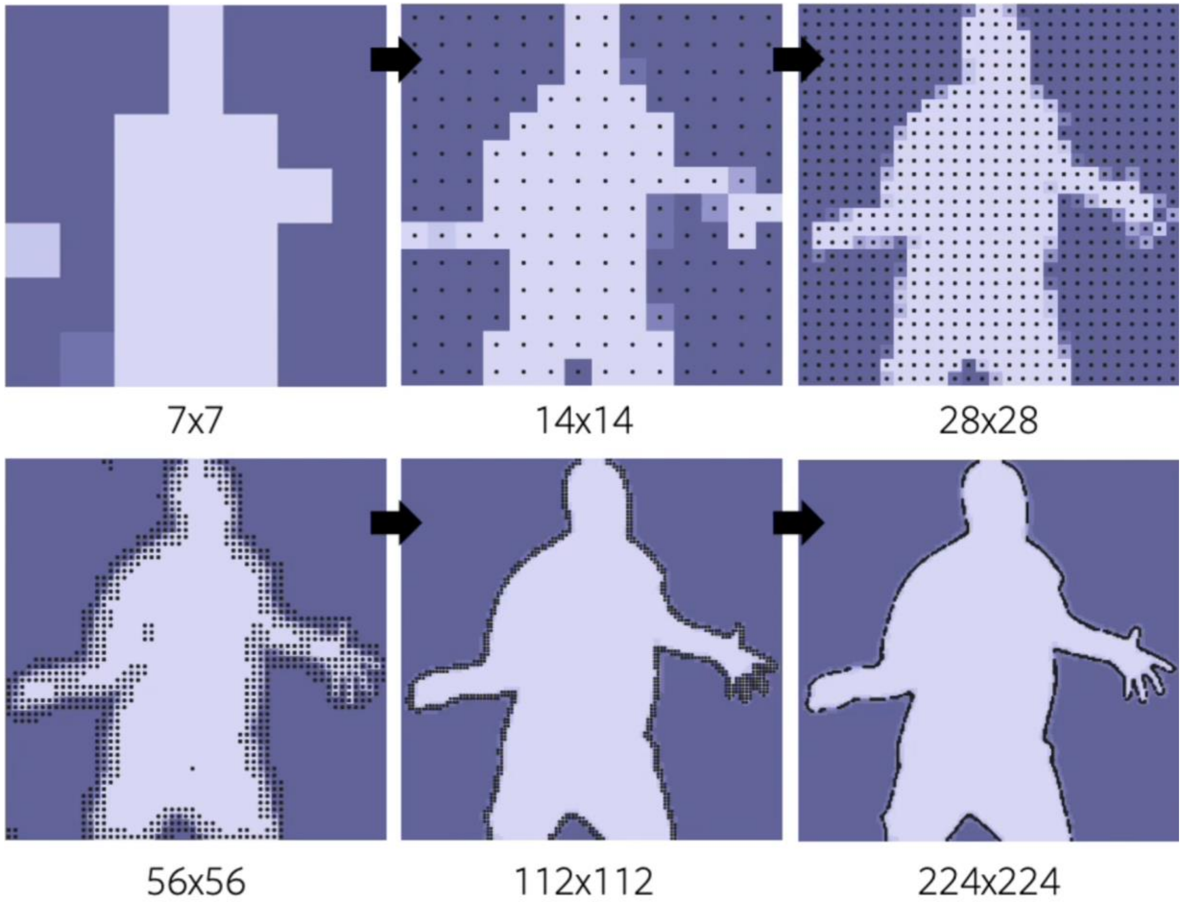

Point-based inference via subdivision

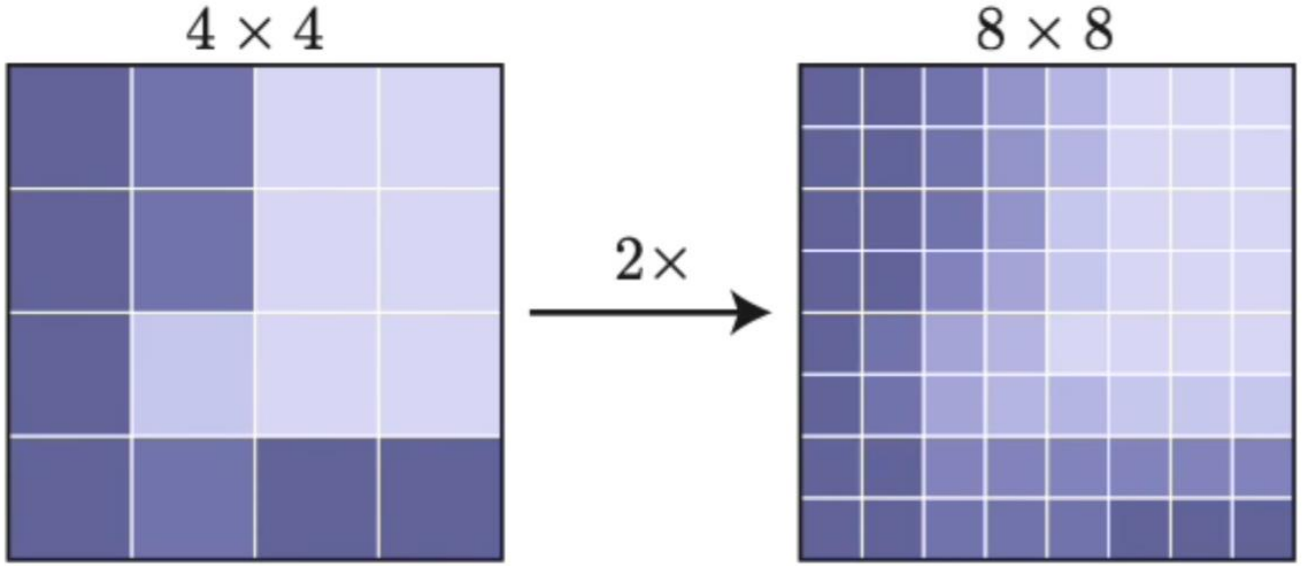

Lower resolutions prediction is
unsampled with bilinear interpolation

Point-based inference via subdivision

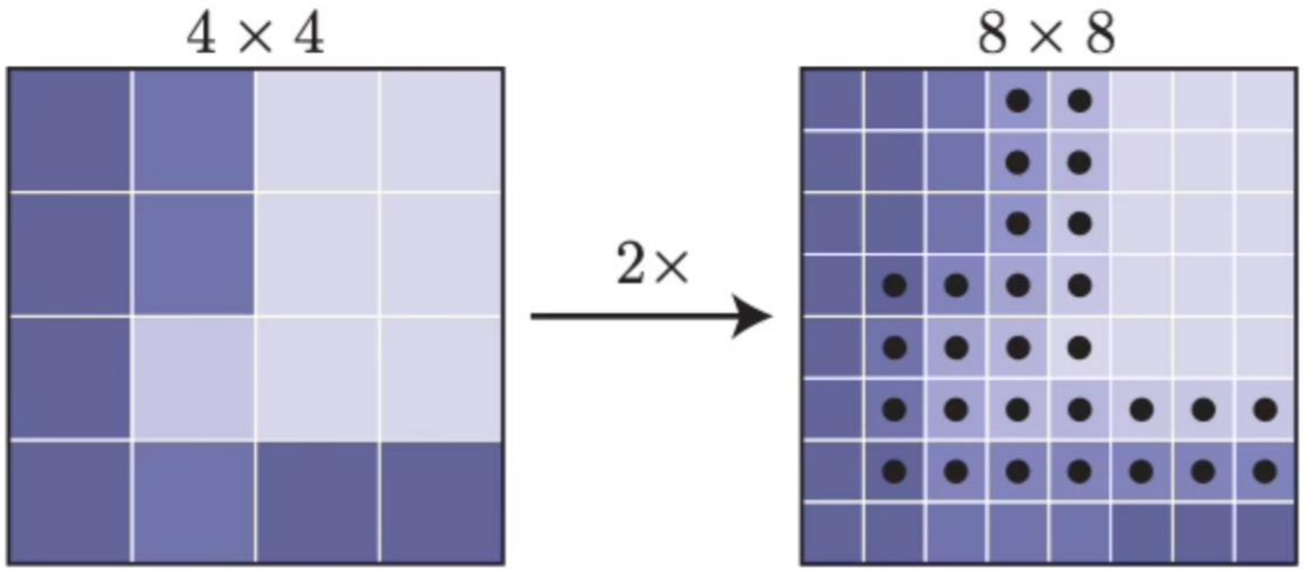

The subset of most uncertain points is selected

Point-based inference via subdivision

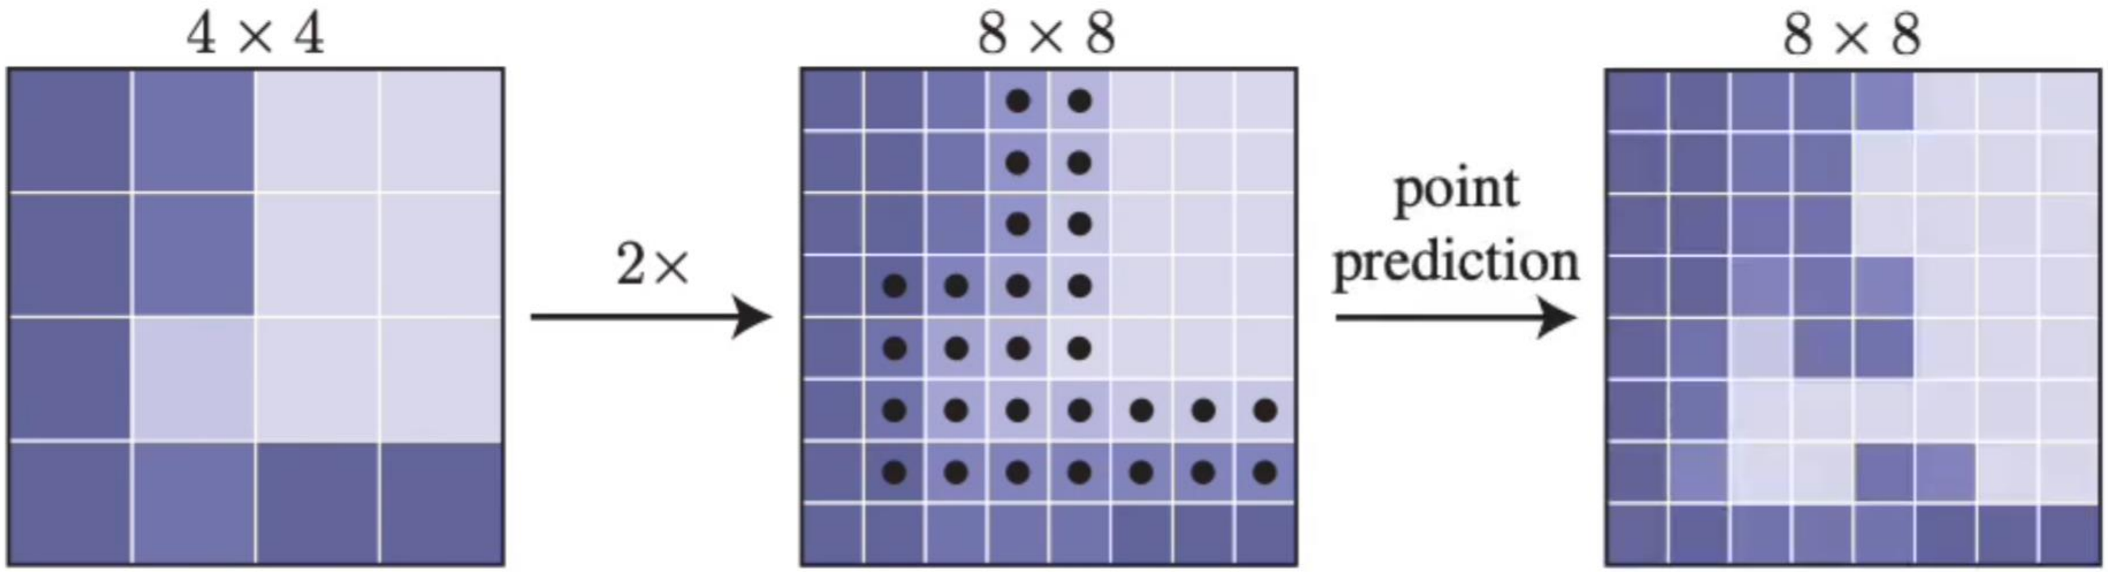

Prediction for each selected points is refined using a lightweight MLP

PointRend architecture for instance segmentation

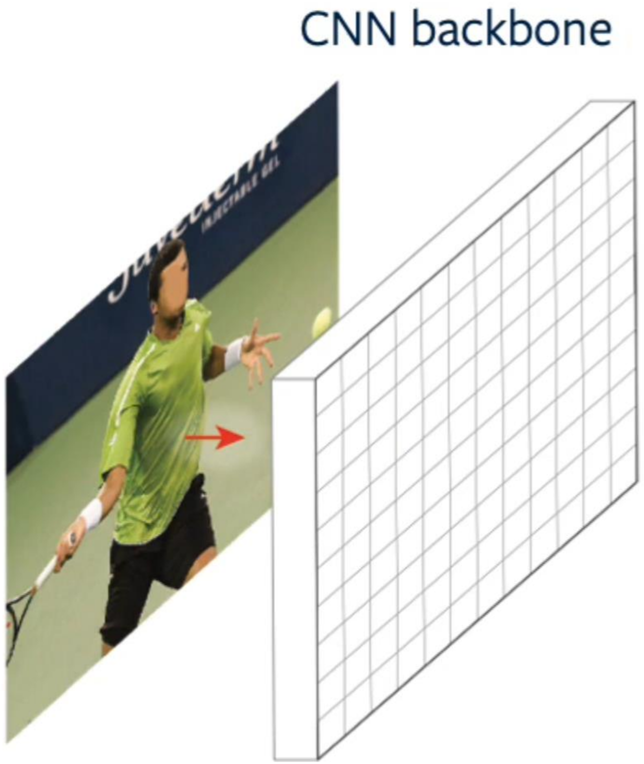

Backbone computes features
that represent the whole image

PointRend architecture for instance segmentation

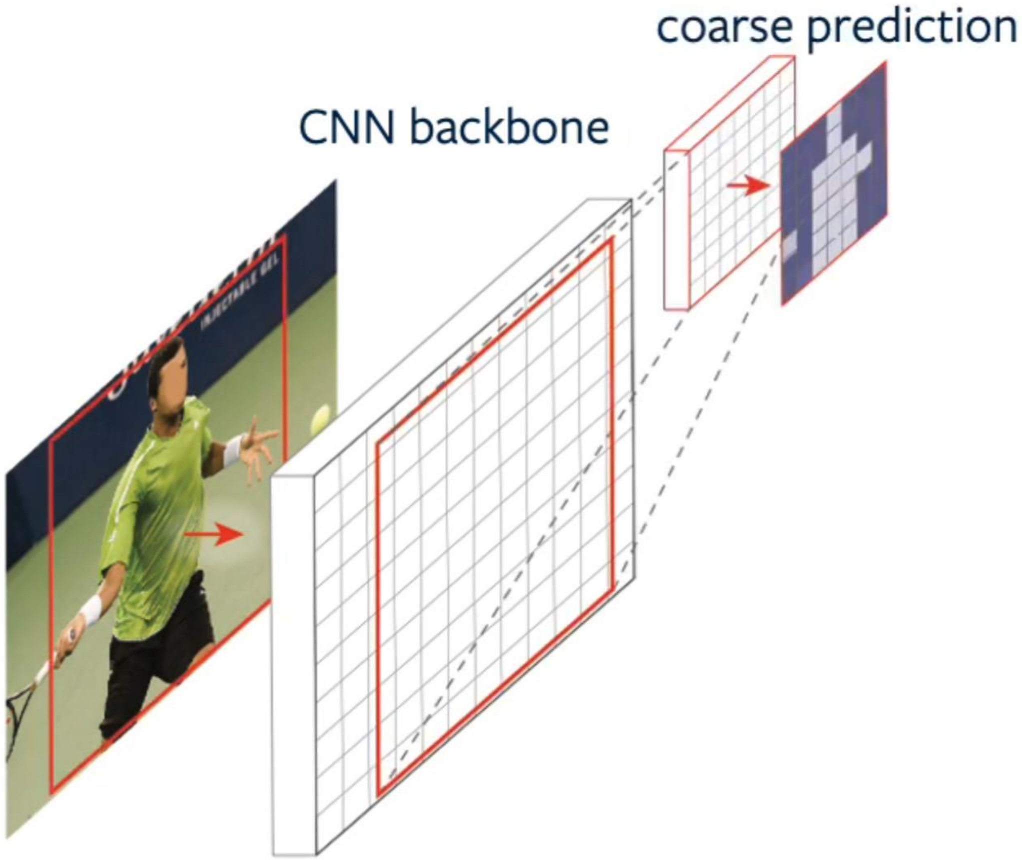

For each bounding box
a small head yields
low-resolution mask prediction

PointRend architecture for instance segmentation

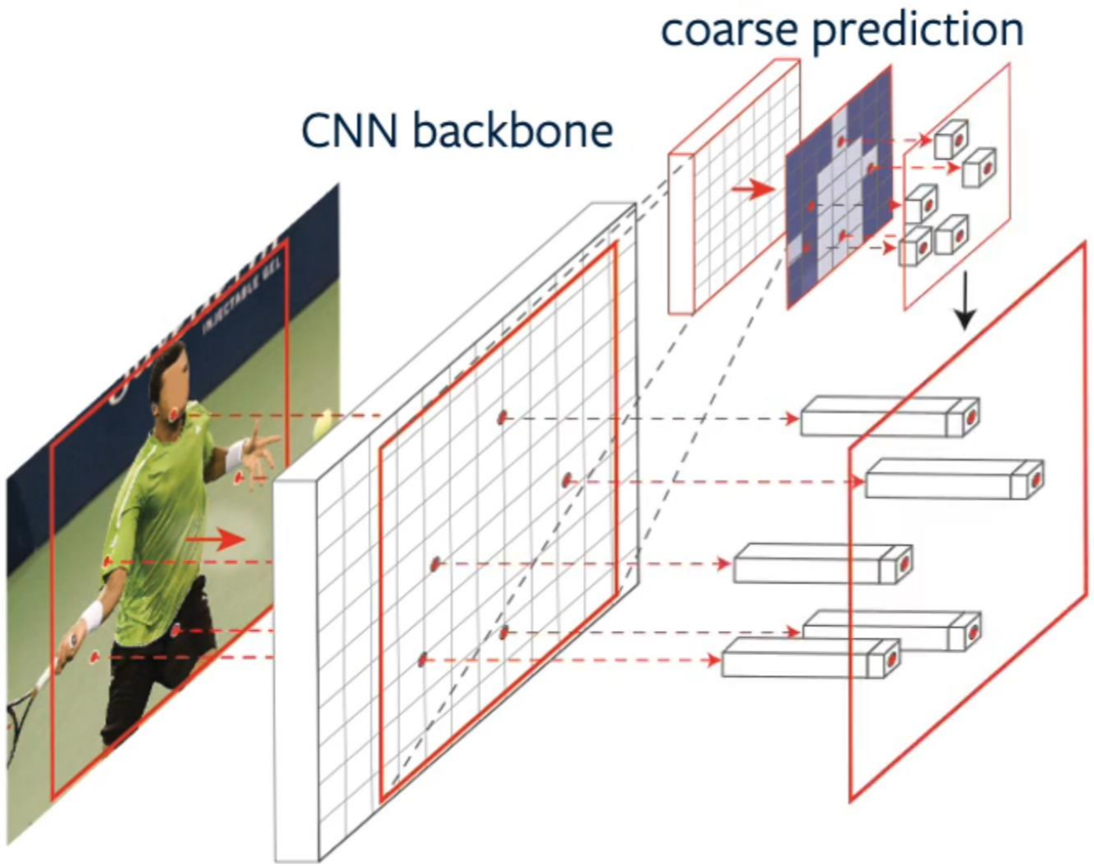

For a subset of points we extract features from the coarse prediction and the backbone features using bilinear interpolation

PointRend architecture for instance segmentation

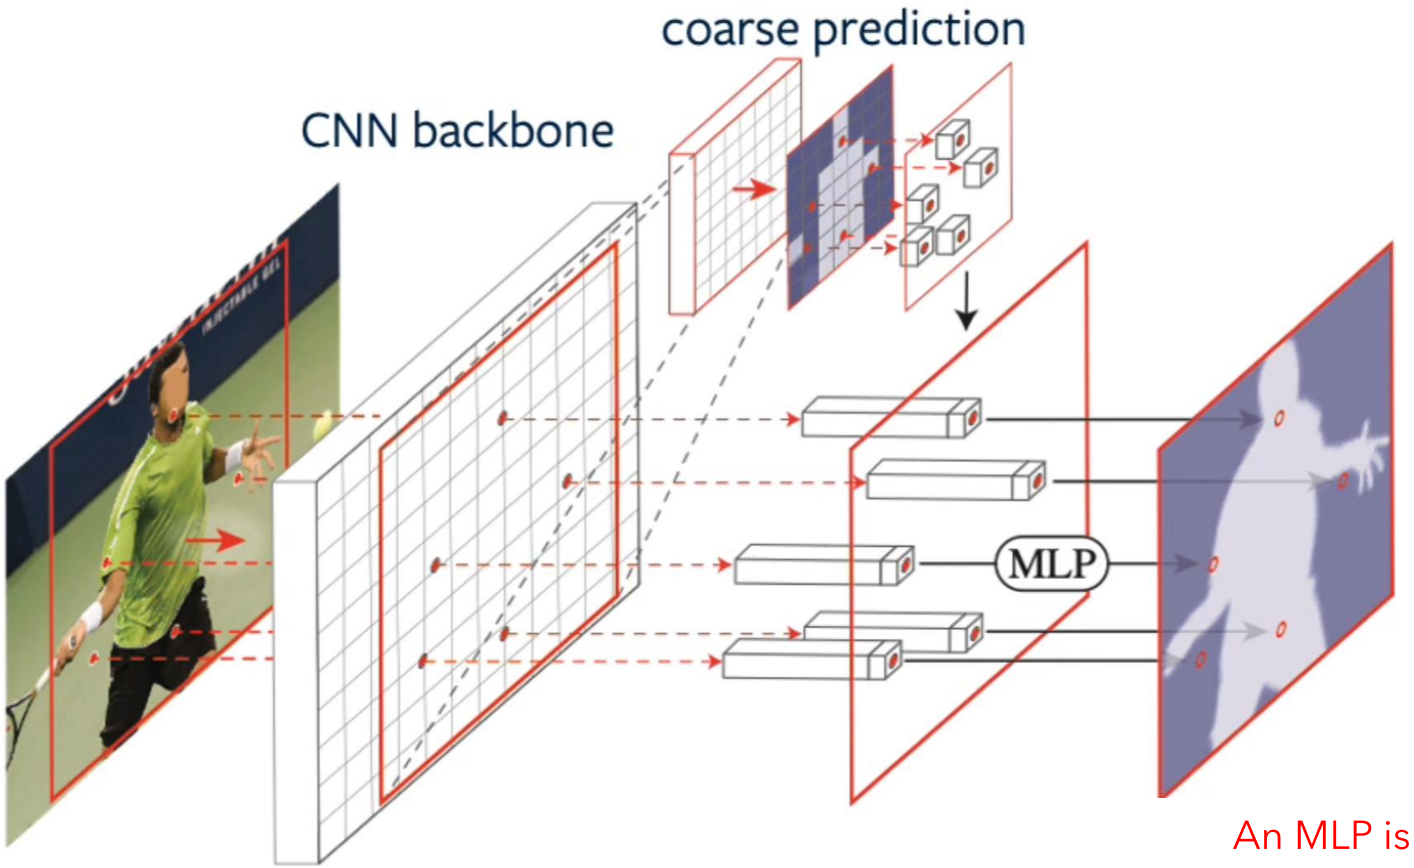

An MLP is used to make prediction
for each point independently

Mask R-CNN with standard head vs. Mask R-CNN with PointRend

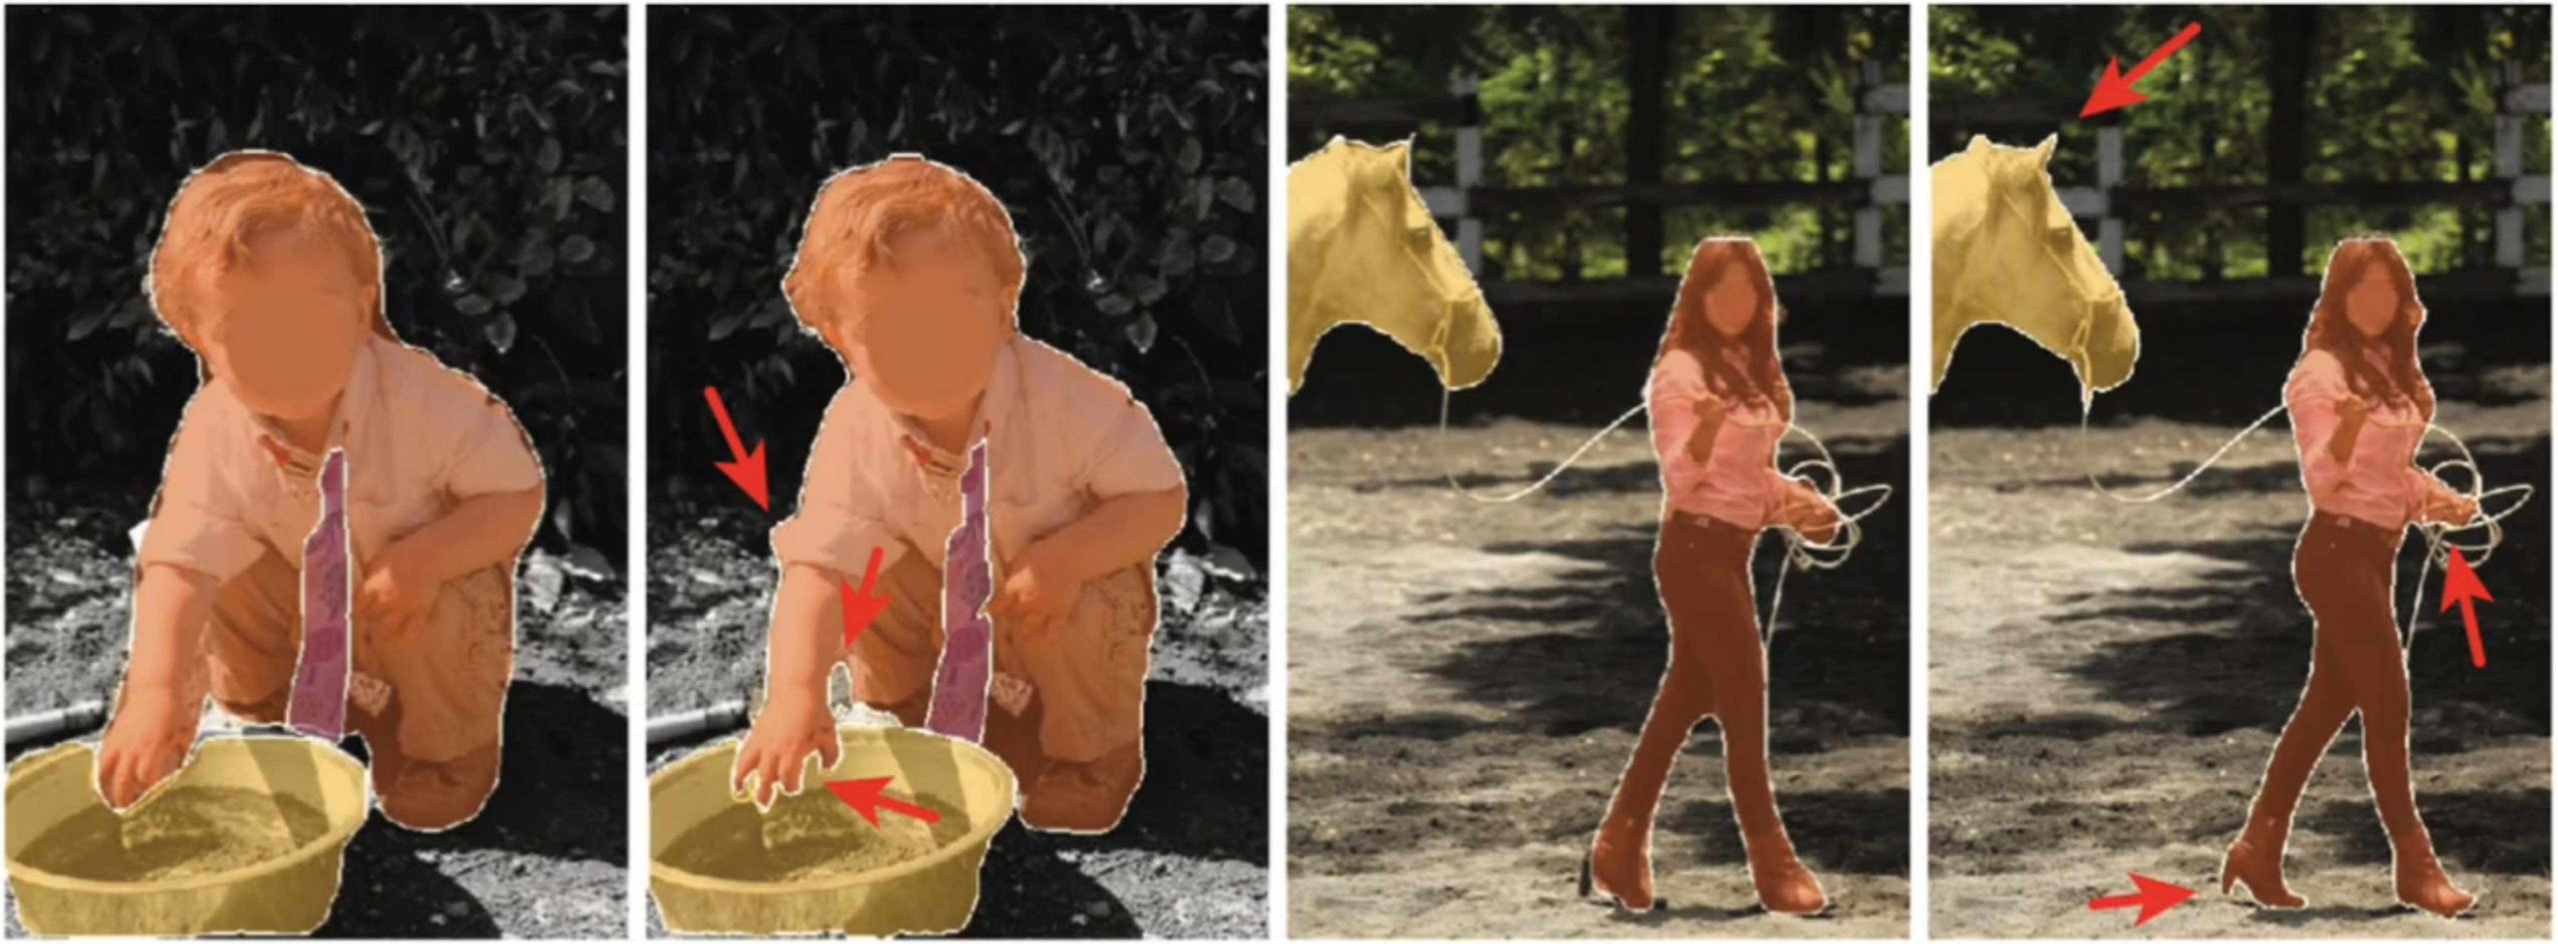

+ PointRend

+ PointRend

Quantitative comparison: instance segmentation

mask head	output resolution	COCO AP	COCO AP*	Cityscapes AP
4× conv	28×28	35.2	37.6	33.0
PointRend	224×224	**36.3** **(+1.1)**	**39.7** **(+2.1)**	**35.8** **(+2.8)**

Mask R-CNN with standard head (4x Conv) vs. Mask R-CNN with PointRend

AP* is COCO mask AP for a COCO-trained model
evaluated against the higher quality LVIS annotation

PointRend for semantic segmentation

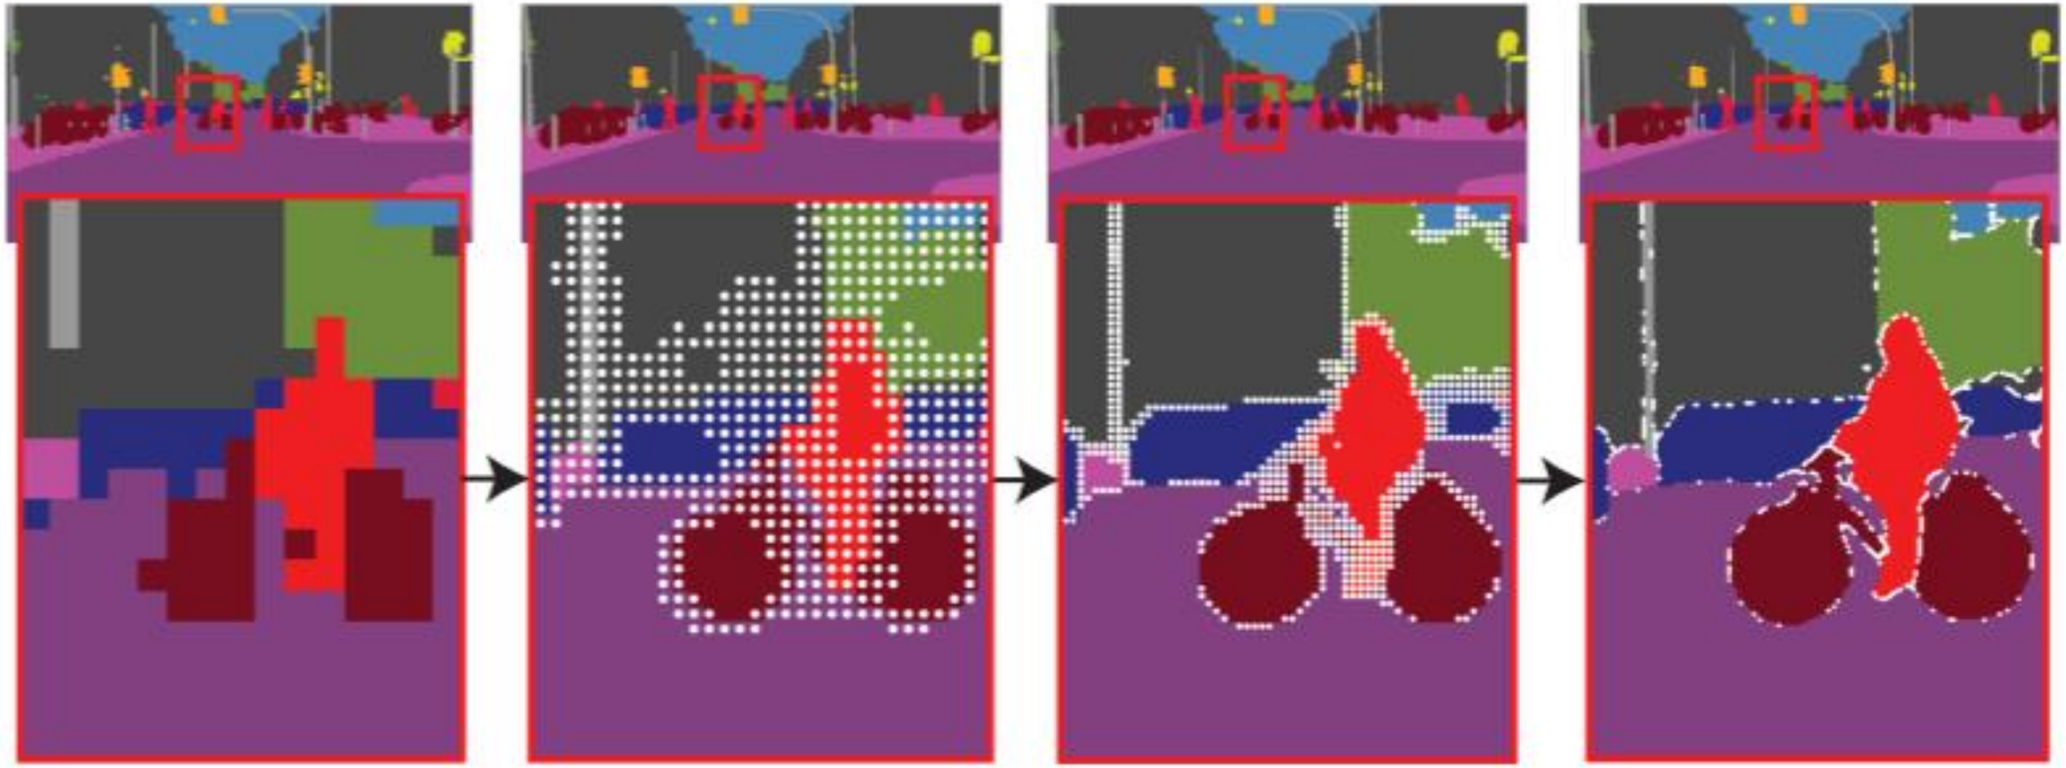

PointRend can be applied on top of
any modern semantic segmentation model

PointRend for semantic segmentation

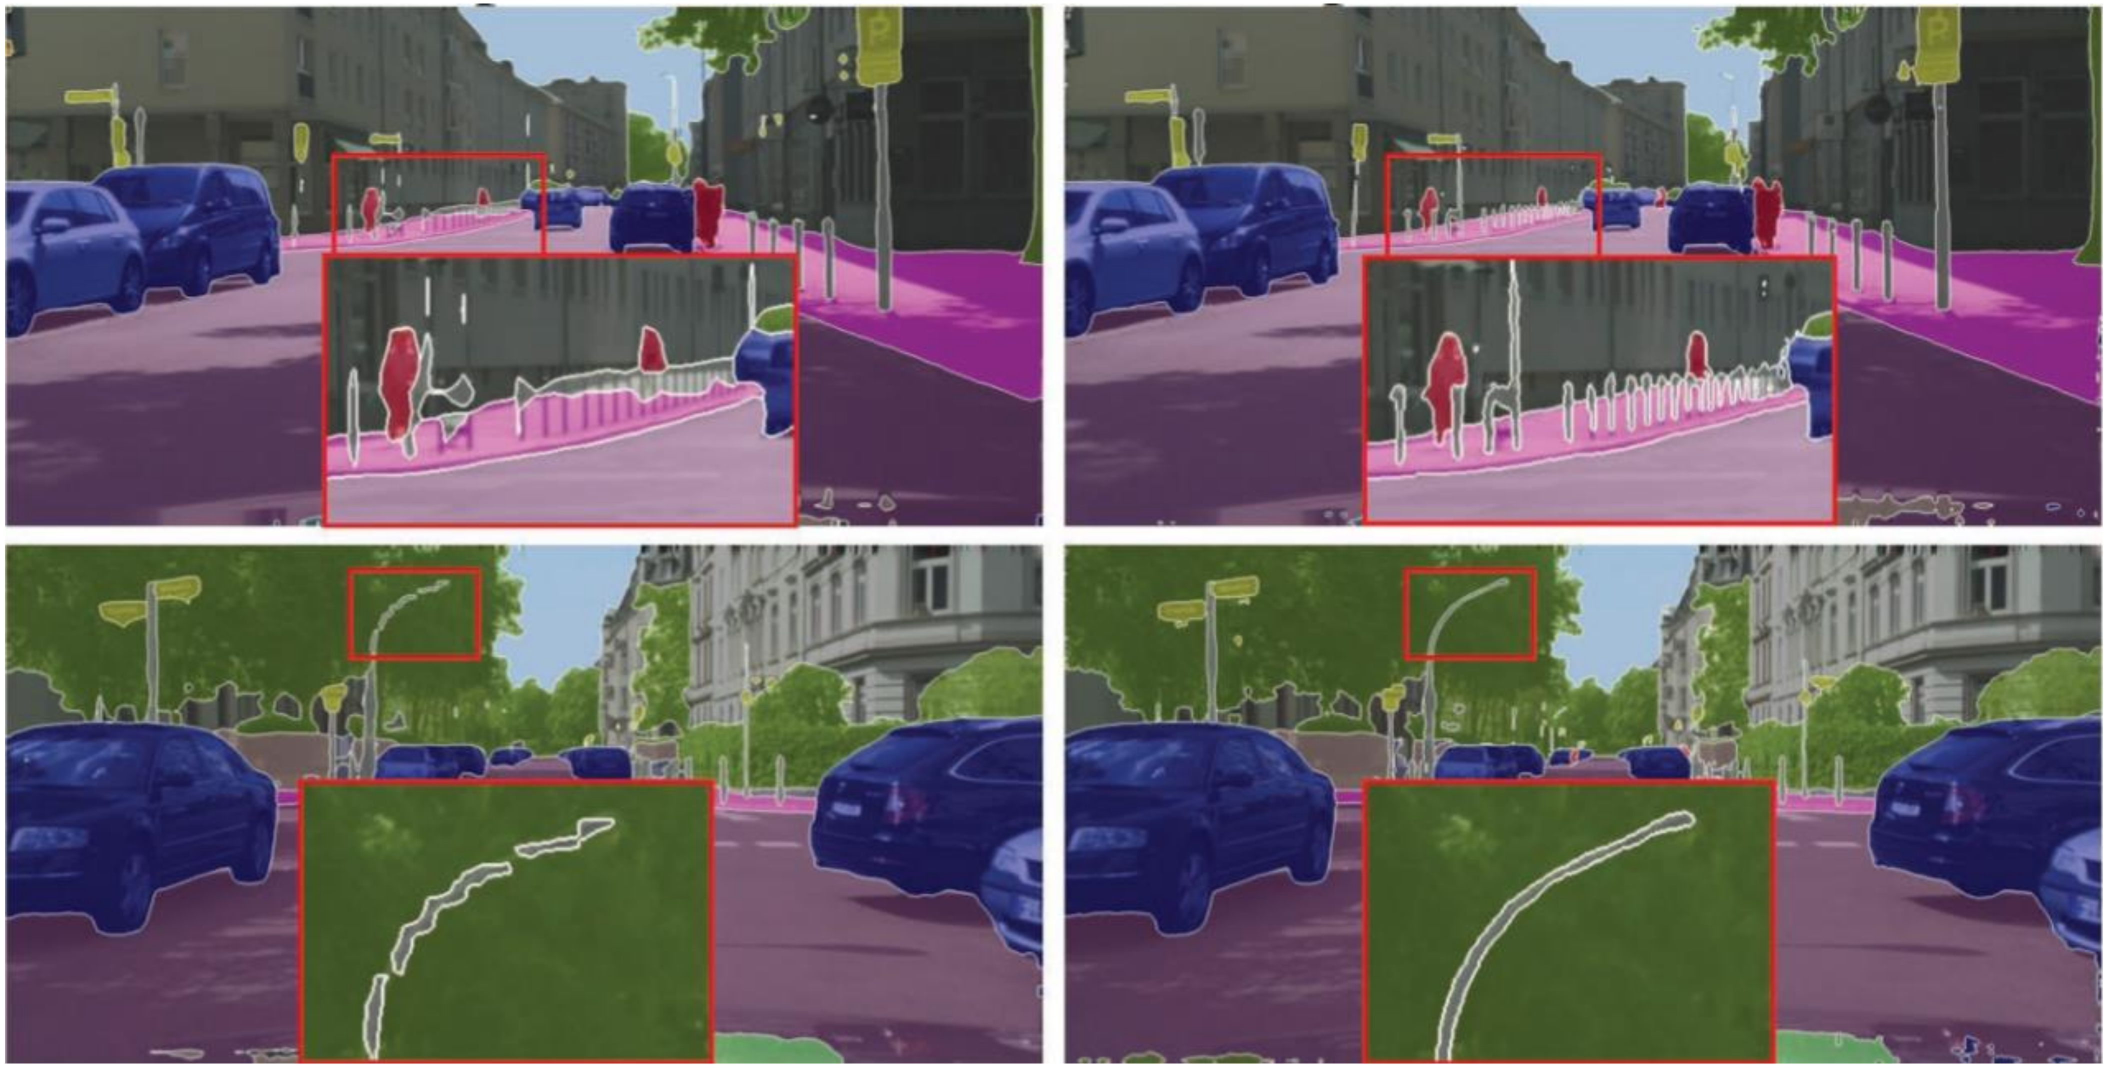

+ PointRend

Quantitative comparison: semantic segmentation

method	output resolution	mIoU
DeeplabV3-OS-16	64×128	77.2
DeeplabV3-OS-8	128×256	77.8 (+0.6)
DeeplabV3-OS-16 + PointRend	1024×2048	**78.4 (+1.2)**

DeeplabV3 vs. DeeplabV3 with PointRend

method	output resolution	mIoU
SemanticFPN P_2-P_5	256×512	77.7
SemanticFPN P_2-P_5 + PointRend	1024×2048	**78.6 (+0.9)**
SemanticFPN P_3-P_5	128×256	77.4
SemanticFPN P_3-P_5 + PointRend	1024×2048	**78.5 (+1.1)**

SemanticFPN vs SemanticFPN with PointRend

Conclusion

- High resolution output with little to no computational overhead

  Higher resolution, more accurate masks

  with fewer model params, less compute time.
- "plug & play" on top of any FCN-based model for segmentation
- Significant quantitative and qualitative improvement

Thank You.

[Paper Review] PointRend: Image Segmentation as Rendering

Su Hyung Choi